

CloudNative Aalborg presents

Introduction to Kubernetes

by Henrik Høegh - Praqma

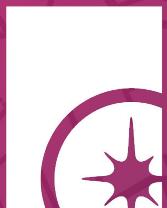
Automating Kubernetes with fluxCD

followed by

Observability using Prometheus, Grafana Loki and Jaeger

by Arne Mejlholm - Sparnord

Hosted by



spar Nord

PRAQMA
Now part of Eficode



CLOUD NATIVE
COMPUTING FOUNDATION

CloudNative Aalborg

#CloudNativeAalborg
#CloudNativeNordics

Join Slack



@ <https://www.cloudnativenoridcs.com/>



Cloud Native Aalborg presents

Introduction to Kubernetes

and

Automating installation of Kubernetes

and

Observability



A word from tonight's sponsor

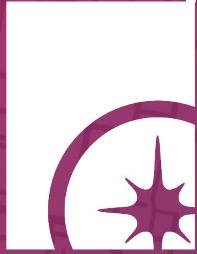


spar Nord



Food

/sponsored by



spar nord



CLOUD NATIVE
COMPUTING FOUNDATION

Do you want to host a Meetup?



Arne Mejlholm
IT Developer
Sparnord



Simon Bengtsson
Systems Engineer
Sparnord



Camilla Beck Larsen
Systems Engineer
Sparnord



Allan Højgaard Jensen
Cloud Platform Architect
Oracle

come talk to us...

Do you want to speak at a Meetup?



Arne Mejlholm
IT Developer
Sparnord

Simon Bengtsson
Systems Engineer
Sparnord

Camilla Beck Larsen
Systems Engineer
Sparnord

Allan Højgaard Jensen
Cloud Platform Architect
Oracle

come talk to us...

Cloud Native Aalborg

We need your feedback...

- what can we do better?
- any ideas for future meetups?
- provide feedback on:  **slack**

Tweet! Tweet!



#CloudNativeAalborg

#CloudNativeNordics

Make some noise about our awesome community... @CloudNativeFdn, #kubernetes, etc.



Cloud Native Nordics Community



CLOUD NATIVE
NORDICS



CLOUD NATIVE
NORDICS

Continue the discussions and meet Cloud Natives from Denmark, Sweden, Norway, and Finland

www.cloudnativenorthernlights.com



slack

Sponsored by
 **Cloud.dk**

Introduction to Kubernetes

/by **Henrik Høegh**
DevOps Consultant Praqma



CLOUD NATIVE
COMPUTING FOUNDATION

PRAQMA

PRAQMA

An Introduction to Kubernetes

(Greek (Κυβερνήτης)-> kee-ver-NEE-tees)



About me

Name: Henrik Høegh

Twitter: @Hoeghh

Doing:

-  CLOUD NATIVE
-  ATASSIAN

Contact: heh@praqma.net





Continuous Delivery and DevOps Conference

**28 OCT 2019
Copenhagen**





HENRIK HØEGH

Storage Operator in Kubernetes - Getting your Feet Wet

DevOps Consultant, Praqla

With Kubernetes becoming the standard way of orchestrating containers, we still need to solve the problem of on-prem storage. Storage and state is one of the most challenging tasks to solve when operating in an on-prem Cloud Native environment. Traditionally companies has been using external storage like NFS, but there are certain problems with this approach. In this talk Henrik will talk about the problems of external storage, how to move storage inside the cluster and automate it with the Rook operator, followed by a live demo. By the end of the talk, the listener will be able to get started with Cloud Native storage in Kubernetes and how to take advantage and flexibility Rook provides.

Twitter: [@HenrikHoegh](#) GitHub: [hoeghh](#) Blog: [Praqla Stories](#)

[Read less](#)



Projects I'm working on



<https://www.praqma.com/products/ask/>

Agenda

Why Kubernetes

Shoulders of giants

Building blocks

Cluster architecture

Why Kubernetes

> History of scaling

- > One process <-> one core
- > One process <-> multi cores
- > One process, many threads <-> multi core

> Future of scaling

- > Monolith <-> multi machines
- > Horizontal scaling
- > Microservices

Agenda

Why Kubernetes

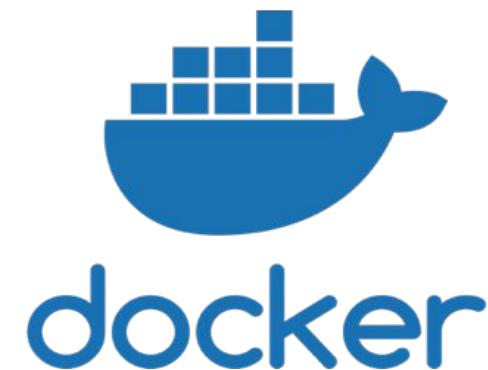
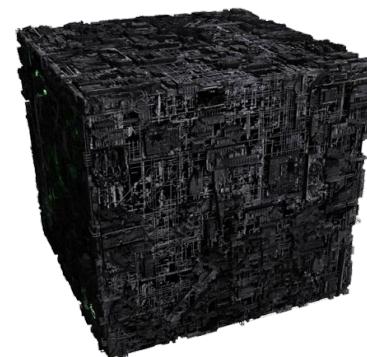
Shoulders of giants

Building blocks

Cluster architecture

Shoulders of giants

Google Borg and Docker.



Agenda

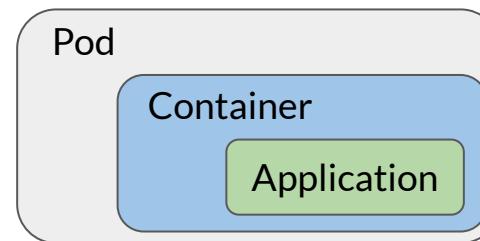
Why Kubernetes

Shoulders of giants

Building blocks

Cluster architecture

Building blocks - Pods



```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
spec:
  containers:
    - name: myapp-container
      image: busybox
      command: ['sh', '-c', 'Application']
```

Agenda

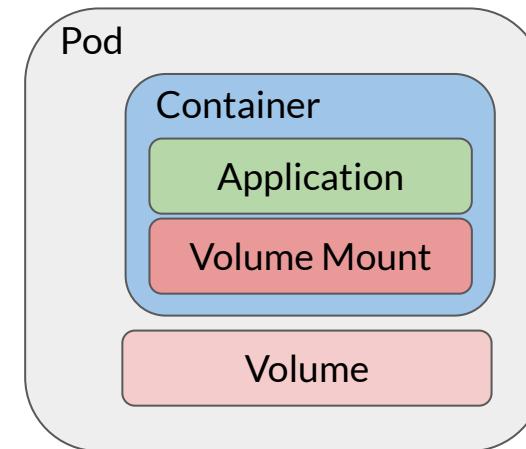
Why Kubernetes

Shoulders of giants

Building blocks

Cluster architecture

Building blocks - Volumes



Agenda

Why Kubernetes

Shoulders of giants

Building blocks

Cluster architecture

Building blocks - Volumes

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
spec:
  containers:
    - name: myapp-container
      image: busybox
      command: ['sh', '-c', 'Application']
        - mountPath: /cache
          name: cache-volume
  volumes:
    - name: cache-volume
      emptyDir: {}
```

Agenda

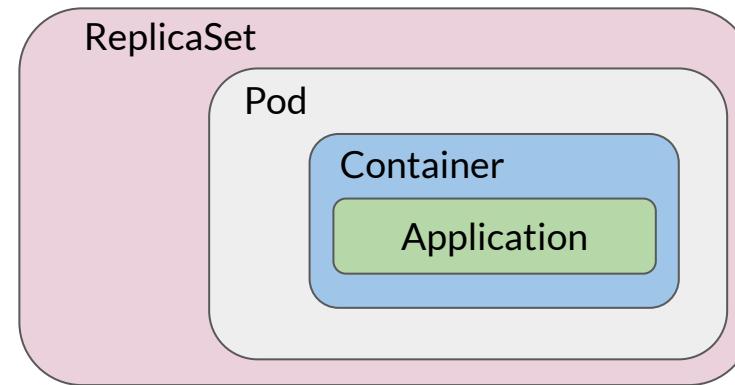
Why Kubernetes

Shoulders of giants

Building blocks

Cluster architecture

Building blocks - ReplicaSets



Agenda

Why Kubernetes

Shoulders of giants

Building blocks

Cluster architecture

Building blocks - ReplicaSets

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: myapp-pod
spec:
  replicas: 3
  selector:
    matchLabels:
      app: myapp
  template:
    metadata:
      labels:
        app: myapp
    spec:
      containers:
        - name: myapp-pod
          image: busybox
```

Agenda

Why Kubernetes

Shoulders of giants

Building blocks

Cluster architecture

Building blocks - ReplicaSets

kubectl get replicaset

NAME	DESIRED	CURRENT	READY	AGE
Myapp-pod	3	0	0	6s

kubectl get replicaset

NAME	DESIRED	CURRENT	READY	AGE
Myapp-pod	3	3	3	18s

`kubectl scale --replicas=5 rs/Myapp-pod`

Agenda

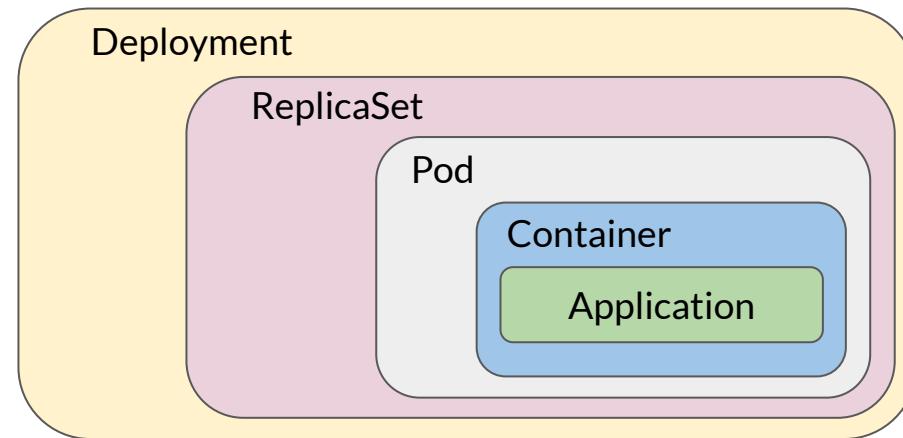
Why Kubernetes

Shoulders of giants

Building blocks

Cluster architecture

Building blocks - Deployment



Agenda

Why Kubernetes

Shoulders of giants

Building blocks

Cluster architecture

Building blocks - Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-pod
  labels:
    app: myapp
spec:
  replicas: 3
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 2 # Max added
      maxUnavailable: 0
  selector:
    matchLabels:
      app: myapp
  template:
    metadata:
      labels:
        app: myapp
    spec:
      containers:
        - name: myapp
          image: busybox
```

Agenda

Why Kubernetes

Shoulders of giants

Building blocks

Cluster architecture

Building blocks - Deployment

```
kubectl get deployments
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
Myapp-pod	3	0	0	0	1s

```
kubectl get deployments
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
Myapp-pod	3	3	3	3	18s

Updating an image version / tag:

```
kubectl --record \
  set image deployment Myapp-pod \
  myapp-container=busybox:1.30.1
```

```
kubectl edit deployment Myapp-pod
```

Agenda

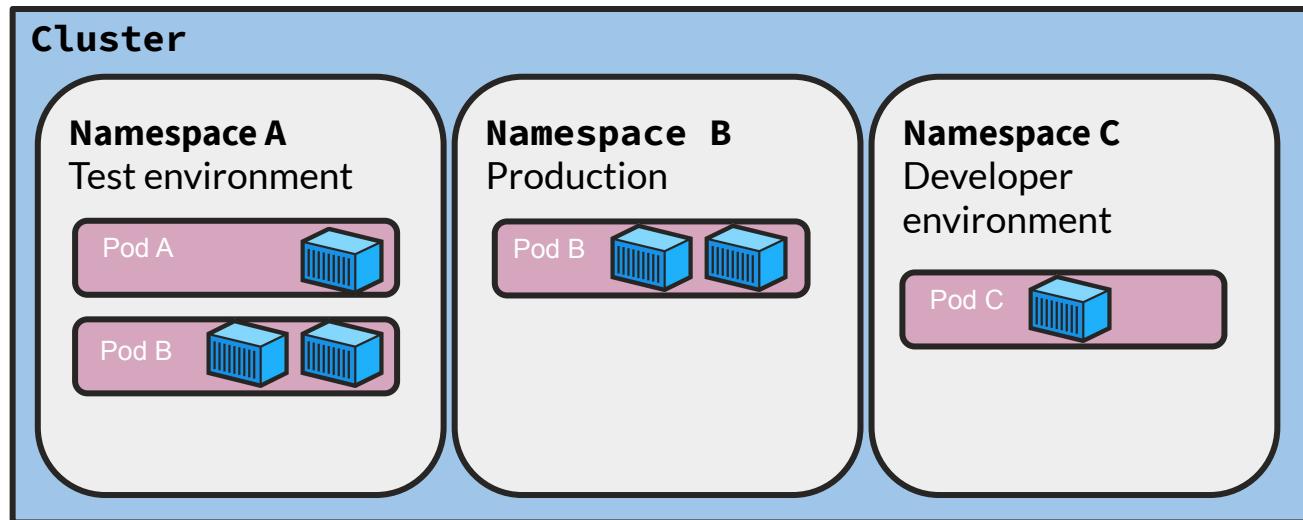
Why Kubernetes

Shoulders of giants

Building blocks

Cluster architecture

Building blocks - Namespaces



Agenda

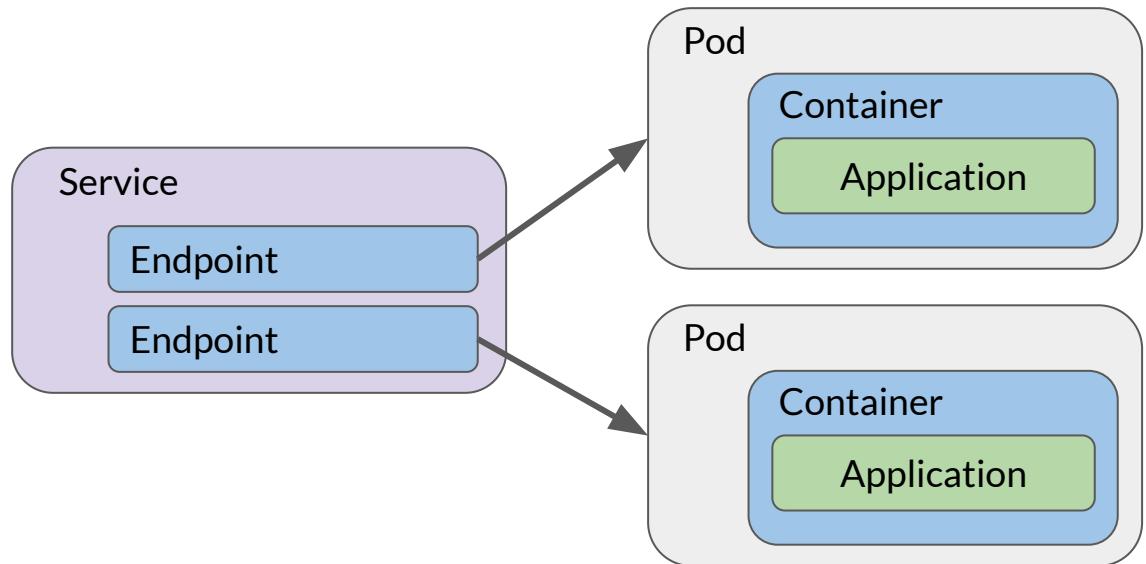
Why Kubernetes

Shoulders of giants

Building blocks

Cluster architecture

Building blocks - Service



Agenda

Why Kubernetes

Shoulders of giants

Building blocks

Cluster architecture

Building blocks - Service

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  selector:
    app: MyApp
  ports:
  - protocol: TCP
    port: 80
    targetPort: 80
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-pod
  labels:
    app: myapp
spec:
  replicas: 3
  selector:
    matchLabels:
      app: myapp
  template:
    metadata:
      labels:
        app: myapp
    spec:
      containers:
      - name: myapp
        image: busybox
```

Agenda

Why Kubernetes

Shoulders of giants

Building blocks

Cluster architecture

Building blocks - Service DNS

[svc-name].[namespace].svc.cluster.local

```
$ curl myapp-svc.prod.svc.cluster.local:80
```

Agenda

Why Kubernetes

Shoulders of giants

Building blocks

Cluster architecture

Building blocks - Service types

- > ClusterIP (“None”?)
- > NodePort
- > ExternalName
- > LoadBalancer

Agenda

—

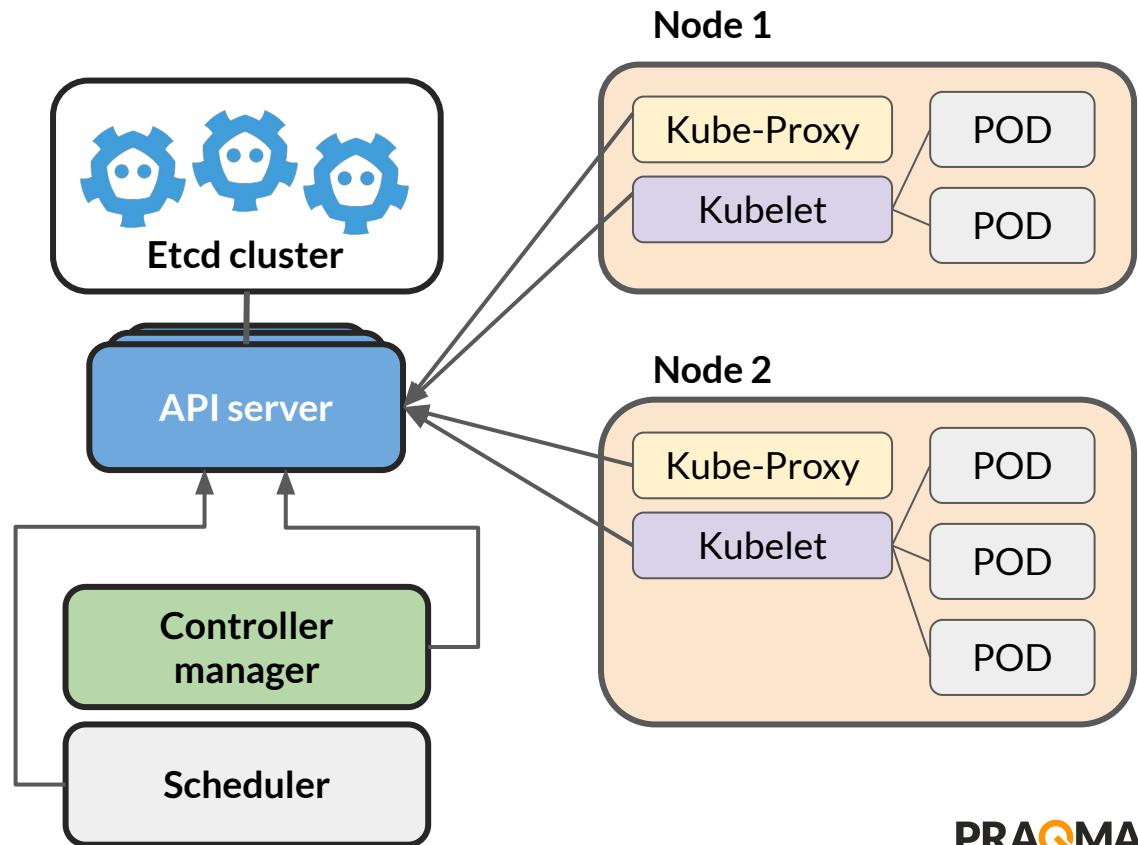
Why Kubernetes

Shoulders of giants

Building blocks

Cluster architecture

Cluster architecture



Agenda

Why Kubernetes

Shoulders of giants

Building blocks

Cluster architecture

Other stuff

Things I didn't cover, that you should look up:

- > Configmaps & Secrets
- > Persistent Volumes & Claims
- > Ingress
- > Other controllers e.g.
 - > Statefulsets
 - > Daemonsets



Thank you

Remember to tweet!

@HenrikHoegh

Automating Kubernetes with fluxCD

/by Arne Mejholm
IT Developer Sparnord



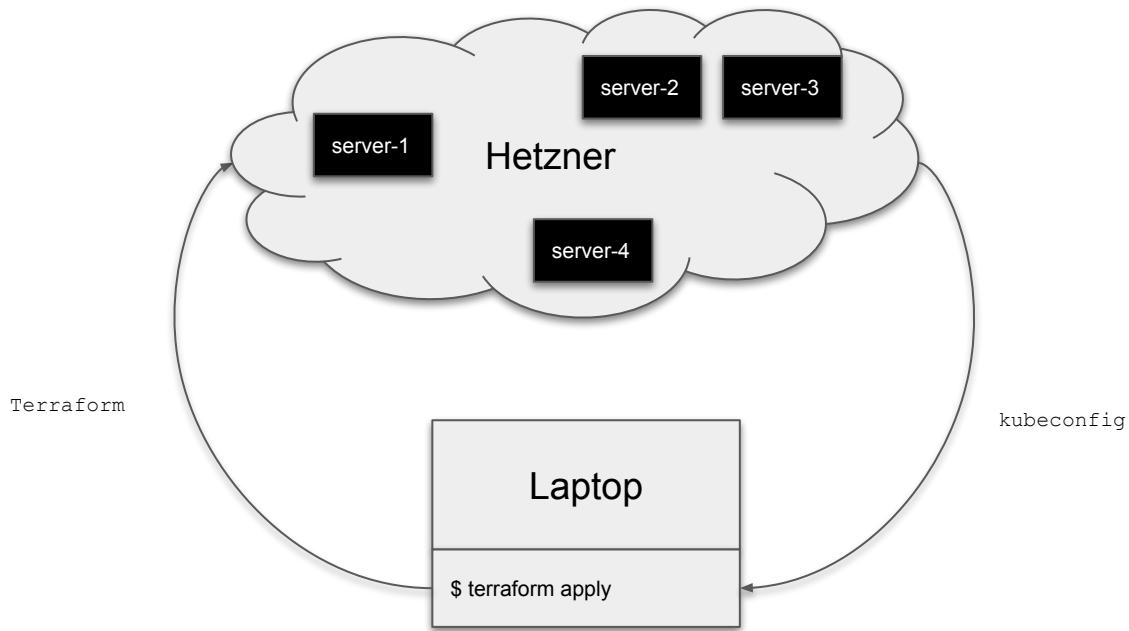
CLOUD NATIVE
COMPUTING FOUNDATION

Infrastructure as code

- Last time we created servers manually
- Sounds more like pets than cattle
- Lets automate that!

Terraform

- Use Infrastructure as Code to provision and manage any cloud, infrastructure, or service



```
##### Variables #####
# Define Hetzner provider
provider "hcloud" {
    token = "<hetzner_cloud_token_goes_here>"
}

# Obtain ssh key data
data "hcloud_ssh_key" "ssh_key" {
    fingerprint = "<ssh_key_goes_here>"
}
```

```
# Create Debian 10 server
resource "hcloud_server" "server-1" {
    name = "server-1"
    image = "debian-10"
    server_type = "cx11"
    ssh_keys  = ["${data.hcloud_ssh_key.ssh_key.id}"]

    connection {
        host = "${hcloud_server.server-1.ipv4_address}"
        type = "ssh"
        user = "root"
        private_key = "${file("~/ssh/id_rsa")}"
    }
}
```

```
provisioner "remote-exec" {
  inline = ["# Connected!"]
}

provisioner "local-exec" {
  command = "k3sup install --ip ${hcloud_server.server-1.ipv4_address} && mv kubeconfig ~/.kube/config"
}

provisioner "local-exec" {
  command = "kubectl create ns apps"
}

provisioner "local-exec" {
  command = "kubectl create secret generic wisdom-service-secret -n apps --from-env-file=secret.env"
}
```

```
##### Variables #####
# Define Hetzner provider
provider "hcloud" {
    token = "<hetzner_cloud_token_goes_here>"
}

# Obtain ssh key data
data "hcloud_ssh_key" "ssh_key" {
    fingerprint = "<ssh_key_goes_here>"
}

# Create Debian 10 server
resource "hcloud_server" "server-1" {
    name = "server-1"
    image = "debian-10"
    server_type = "cx11"
    ssh_keys = ["${data.hcloud_ssh_key.ssh_key.id}"]

    connection {
        host = "${hcloud_server.server-1.ipv4_address}"
        type = "ssh"
        user = "root"
        private_key = "${file("~/ssh/id_rsa")}"
    }

    provisioner "remote-exec" {
        inline = ["# Connected!"]
    }

    provisioner "local-exec" {
        command = "k3sup install --ip ${hcloud_server.server-1.ipv4_address} && mv kubeconfig ~/.kube/config"
    }

    provisioner "local-exec" {
        command = "kubectl create ns apps"
    }

    provisioner "local-exec" {
        command = "kubectl create secret generic wisdom-service-secret -n apps --from-env-file=secret.env"
    }
}
```

GitOps

GitOps is a way to do Kubernetes cluster management and application delivery.

It works by using Git as a single source of truth for declarative infrastructure and applications.

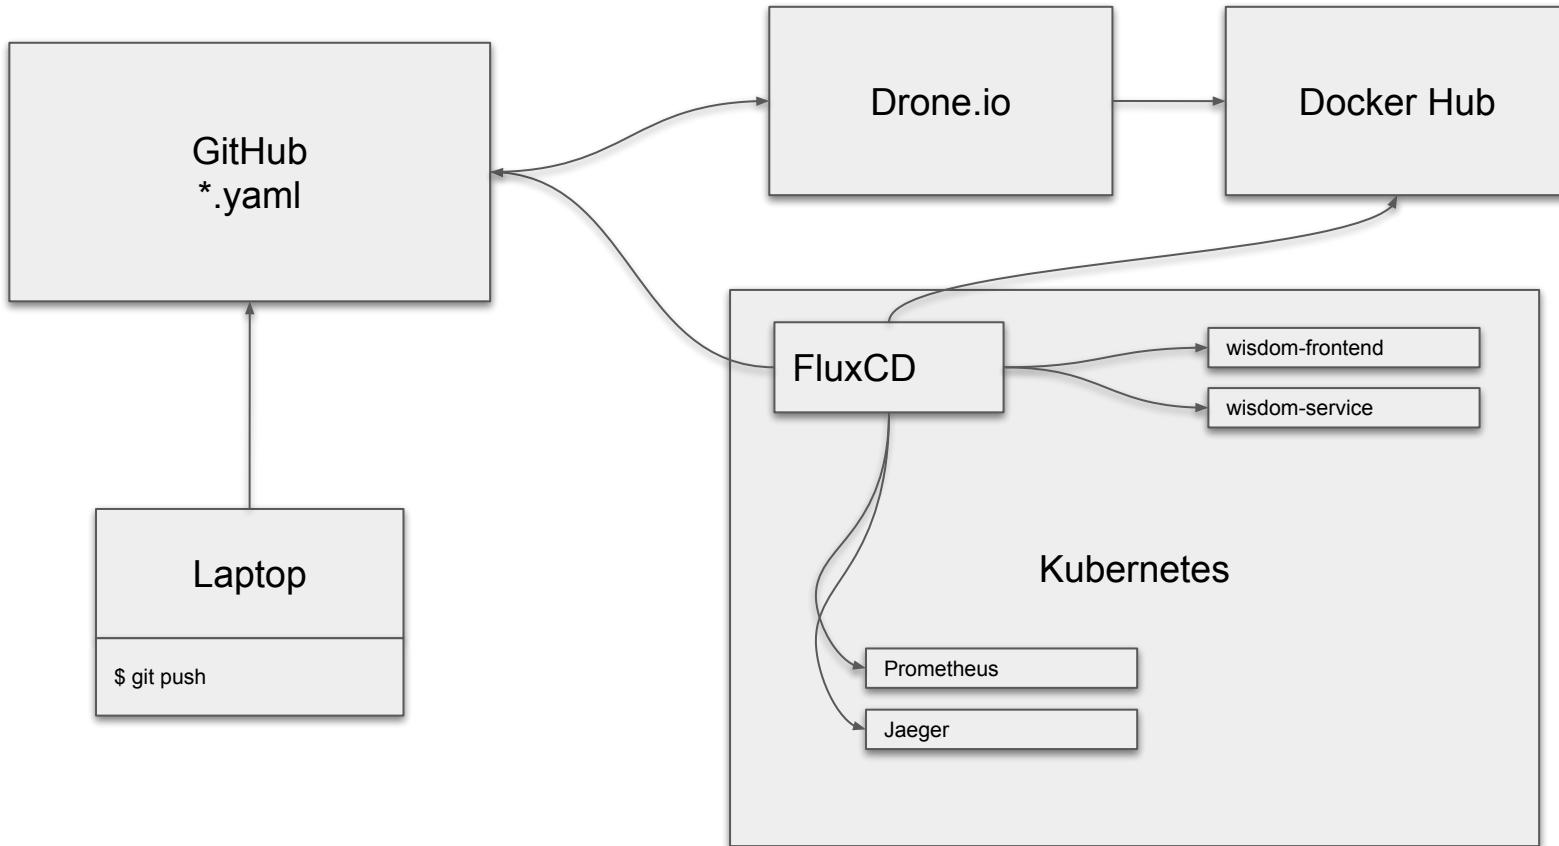
FluxCD

- The GitOps operator for Kubernetes

```
fluxctl install --git-user=mejlholm \
--git-email=mejlholm@users.noreply.github.com \
--git-url=git@github.com:Cloud-Native-Aalborg/Meetup-2 \
--git-paths=k8s,wisdom-frontend/deploy,wisdom-service/deploy
--namespace=flux \
| kubectl apply -f -
```

Drone.io

- Drone is a self-service Continuous Delivery platform for busy development teams.



Lets see that in practice

- Demo time

Observability

with Prometheus, Grafana Loki and Jaeger

/by Arne Mejholm
IT Developer Sparnord



CLOUD NATIVE
COMPUTING FOUNDATION

How are our applications doing?

- What do we do if our applications are misbehaving?
- How many times are our application being used?

Lets install some observability tools:

- Jaeger
- Prometheus
- Grafana Loki

Jaeger



JAEGER

Jaeger: open source, **end-to-end distributed tracing**

Monitor and troubleshoot transactions in complex
distributed systems

Prometheus



Power your **metrics** and alerting with a leading open-source monitoring solution.

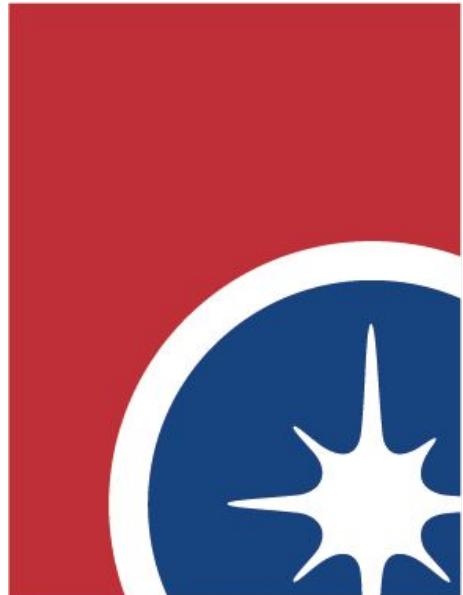
Grafana Loki

Loki is a horizontally-scalable, highly-available, multi-tenant **log aggregation system** inspired by Prometheus.

Lets see that in practice

- Demo time

A big thank you to this evening's host



sparNord

Networking

/by You

Cloudnatives in Aalborg



CLOUD NATIVE
COMPUTING FOUNDATION

CloudNative Aalborg

#CloudNativeAalborg
#CloudNativeNordics

Join Slack



@ <https://www.cloudnativenoridcs.com/>

