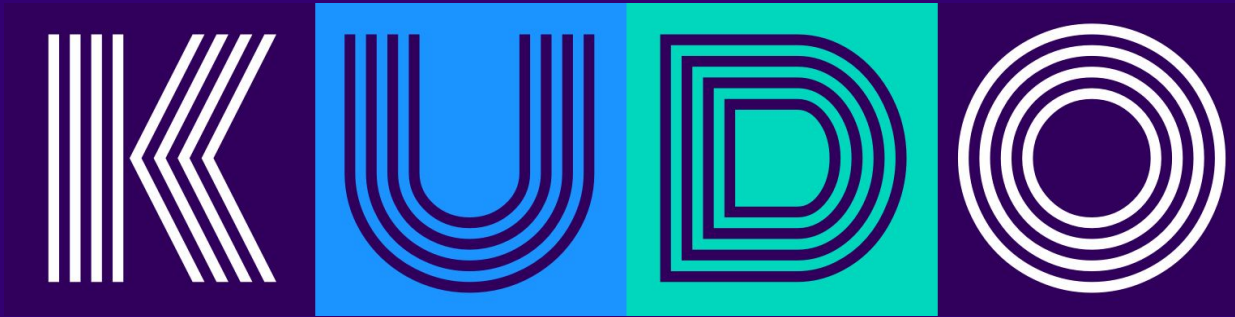


---

# Introducing



---

Kubernetes Operators - The Easy Way

Nick Jones, D2iQ

\$ whoami

## Nick Jones


- Community Engineering Lead @ D2iQ
- Building stuff with open source software for ~20 years
- Ops, Dev and Dev/Ops
- Relatively new to Kubernetes
- ... but not new to Ops
- Likes cats




@yankcrime

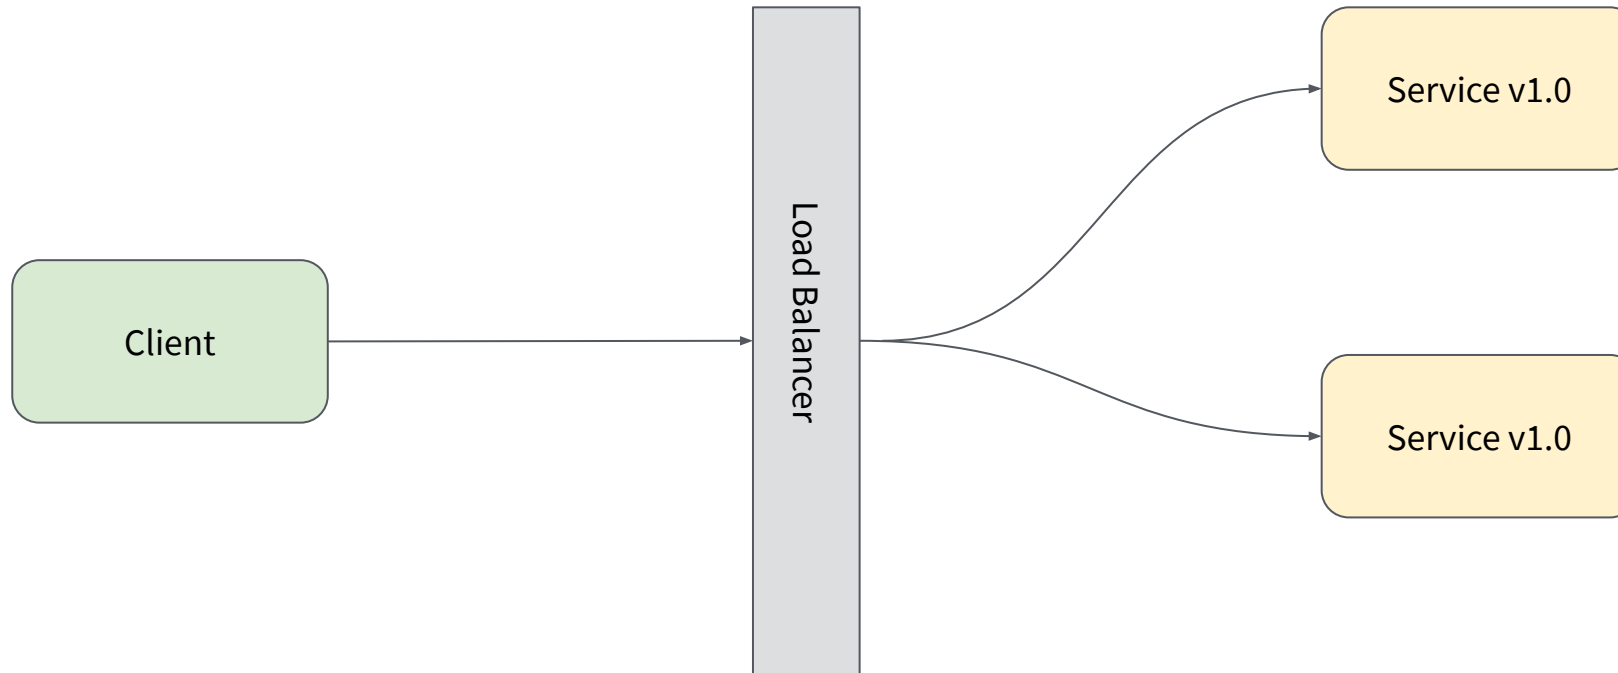


yankcrime



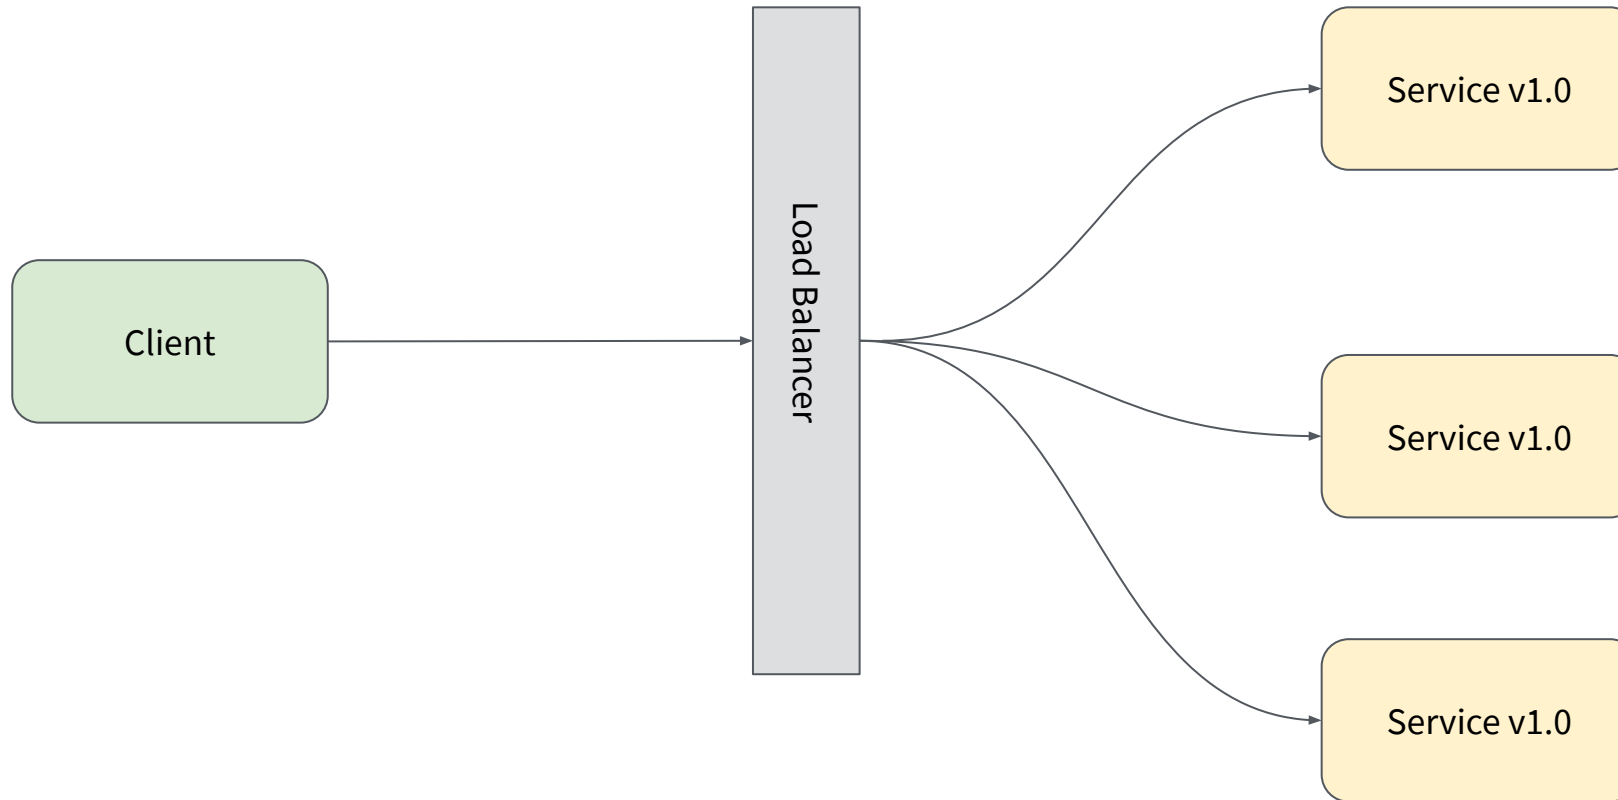
- Stateful vs. Stateless
- Kubernetes StatefulSets
- Kubernetes Operators
- 
  - Background
  - Concepts
  - Demo 🙌
  - Future
  - Getting involved

# Stateless Applications



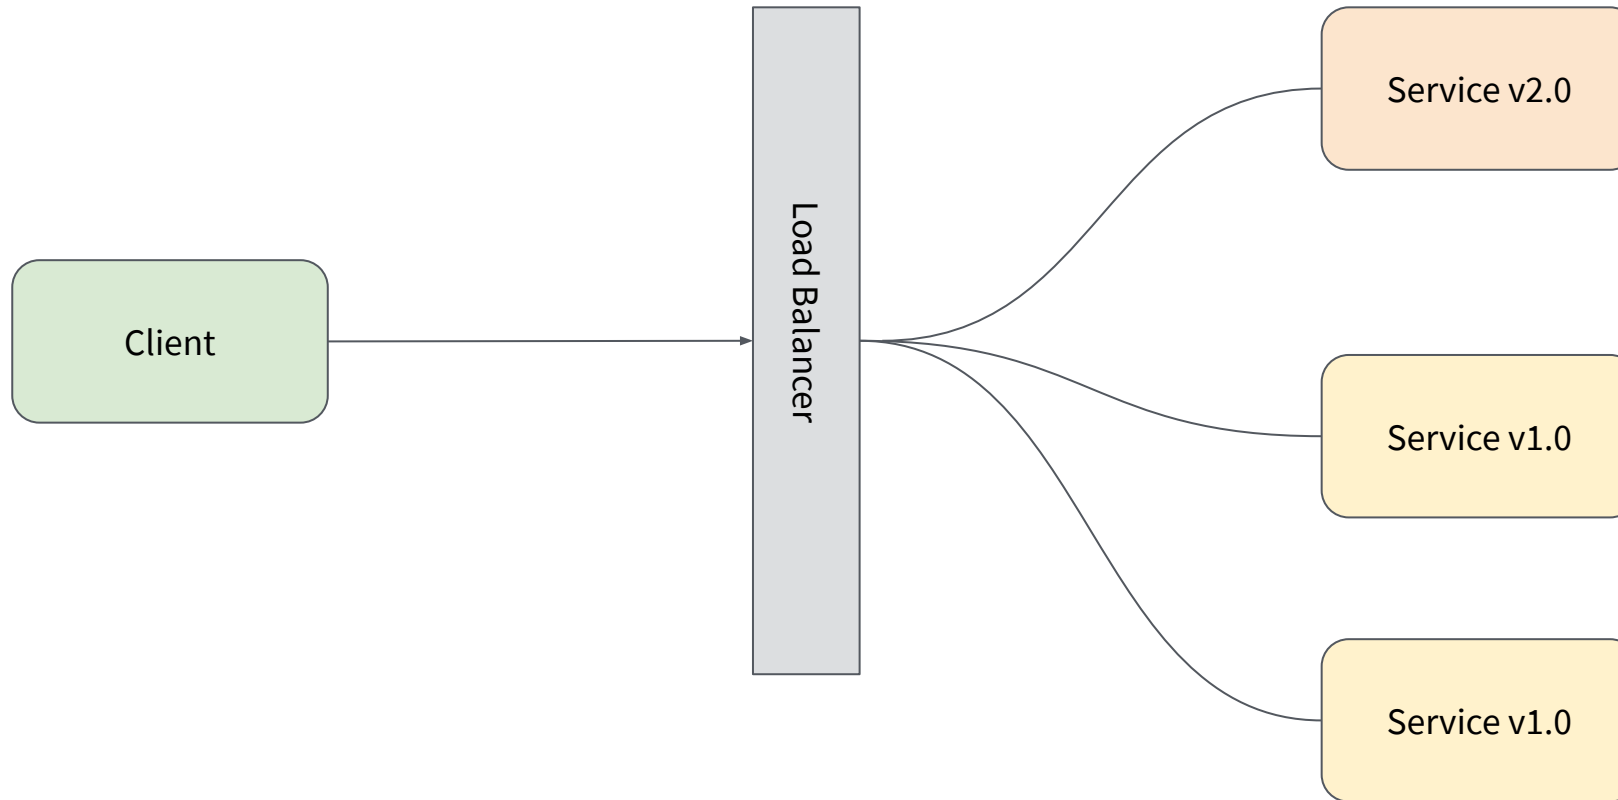
- No state persisted

# Stateless Applications



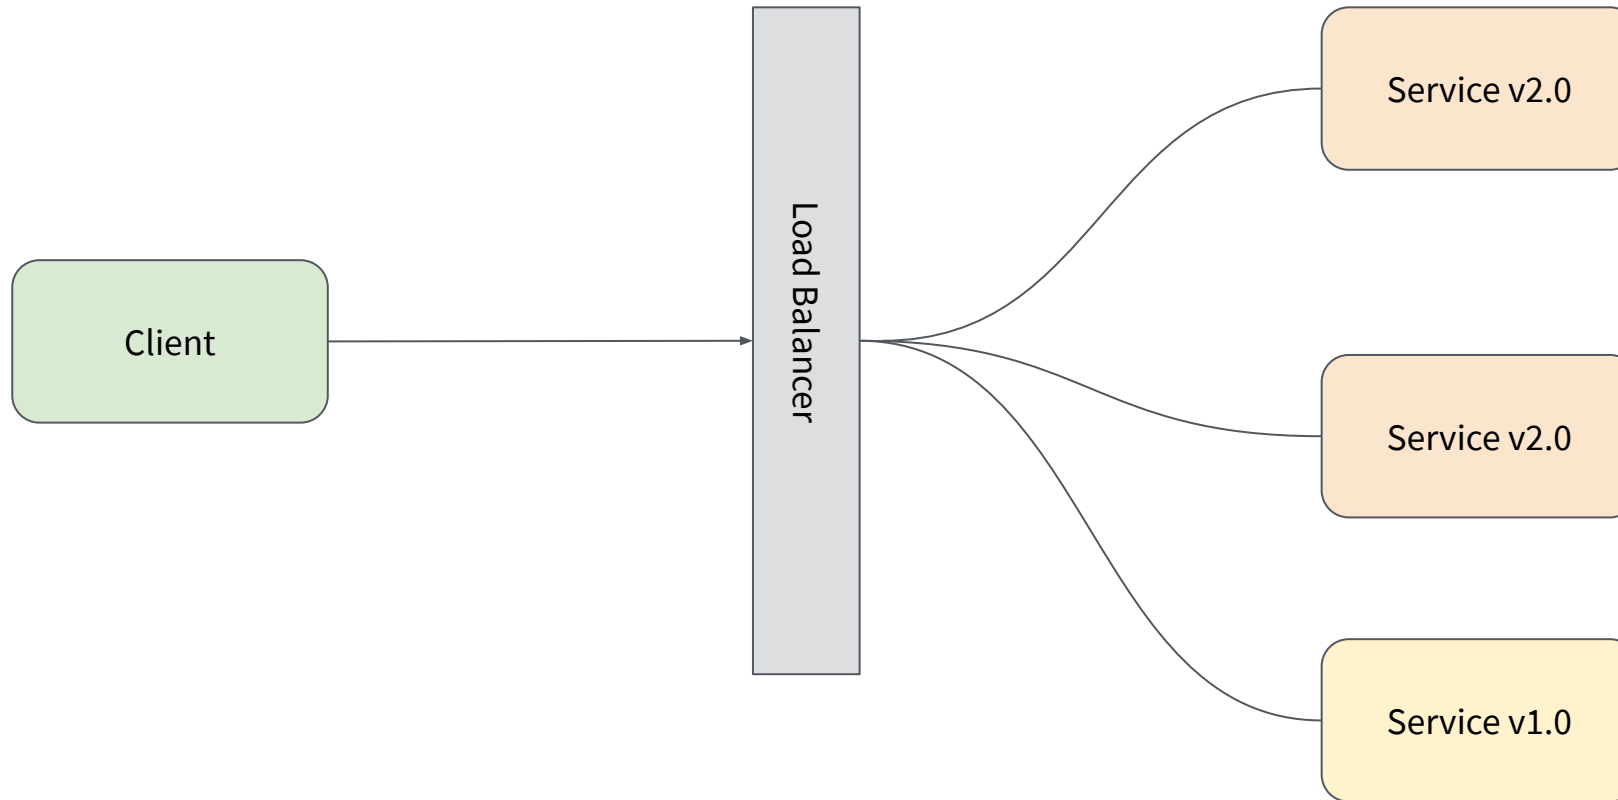
- No state persisted
- Easy to scale up / down

# Stateless Applications



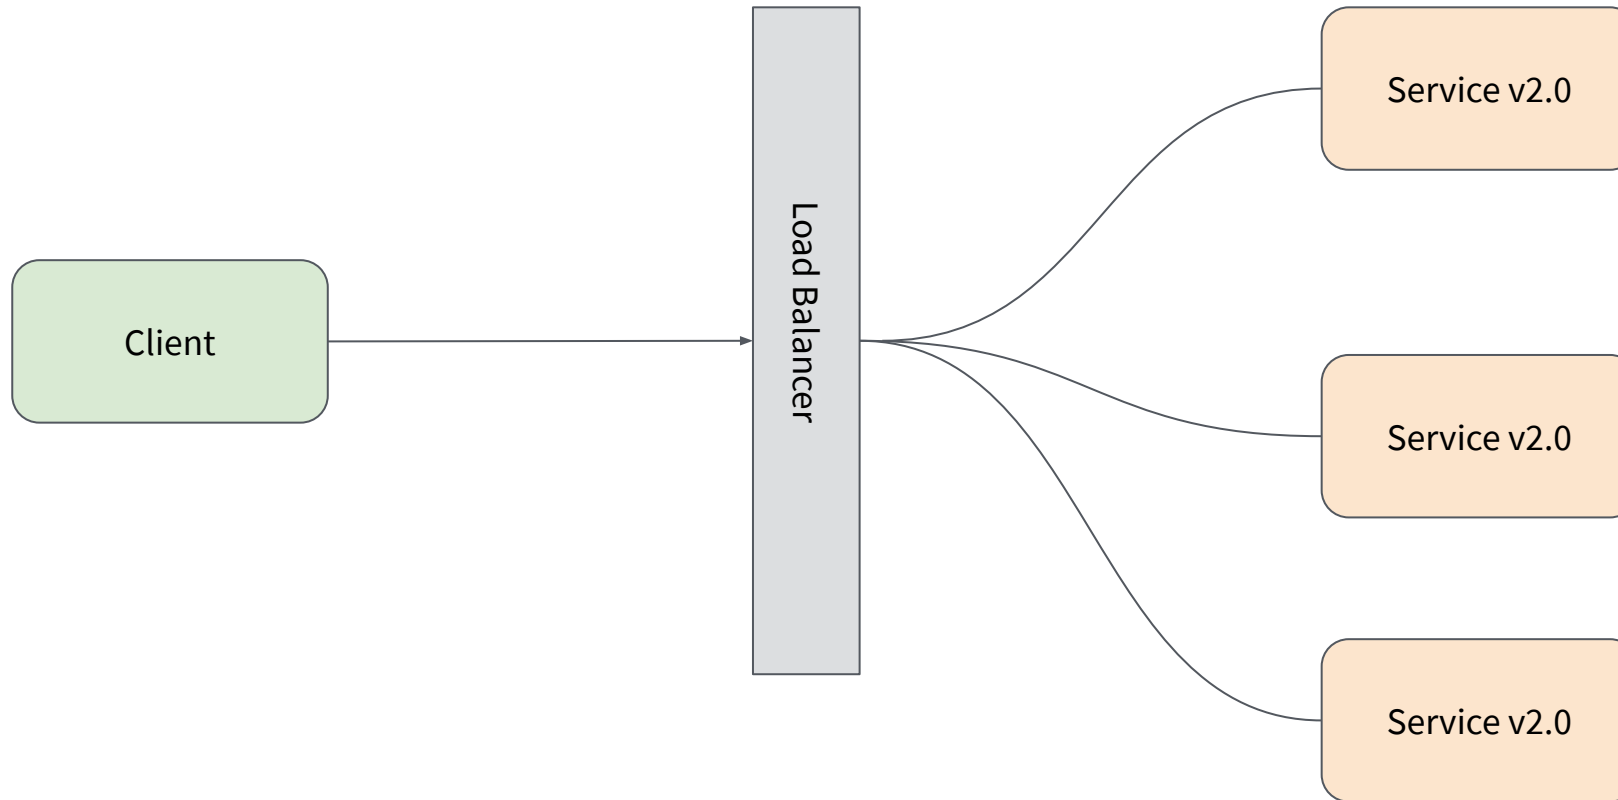
- No state persisted
- Easy to scale up / down

# Stateless Applications



- No state persisted
- Easy to scale up / down

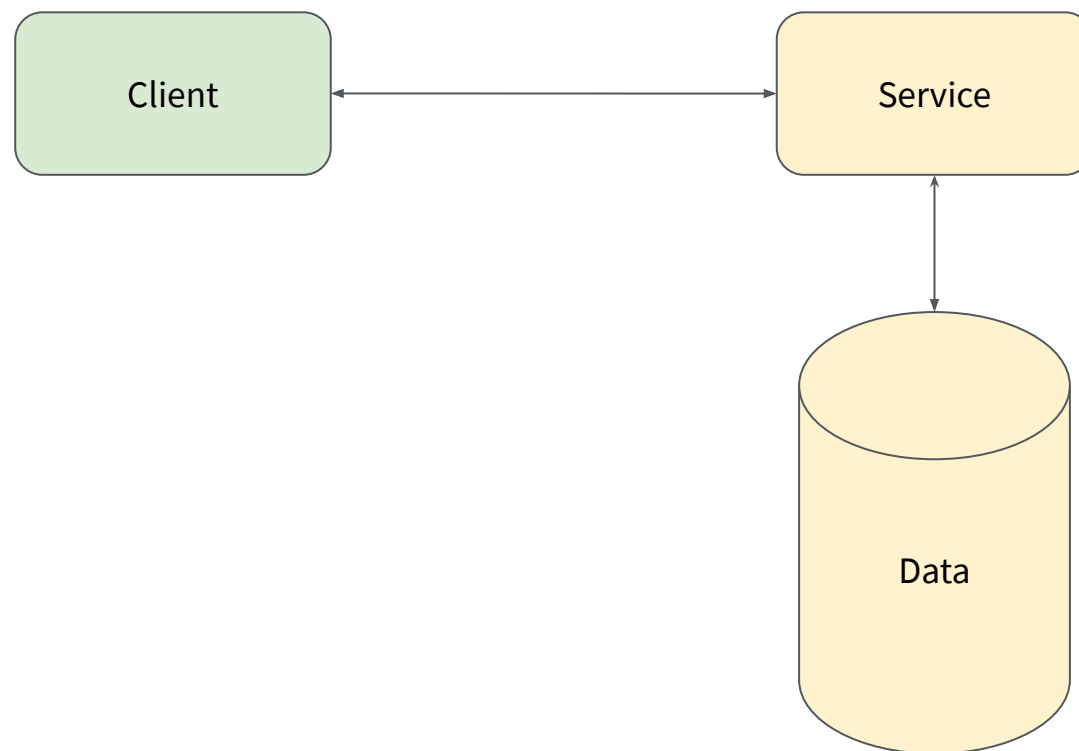
# Stateless Applications



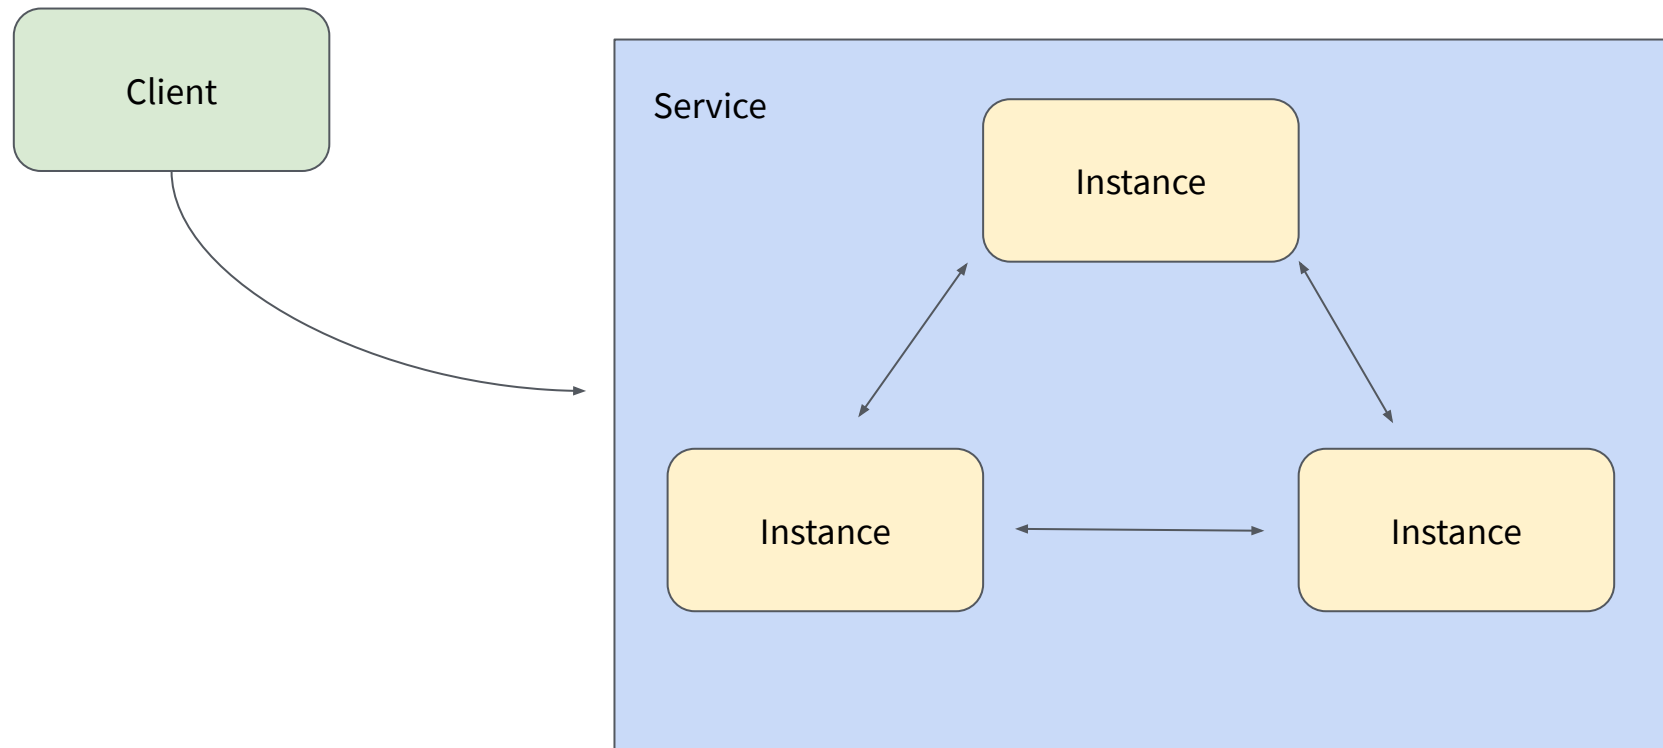
- No state persisted
- Easy to scale up/down
- Straightforward to upgrade



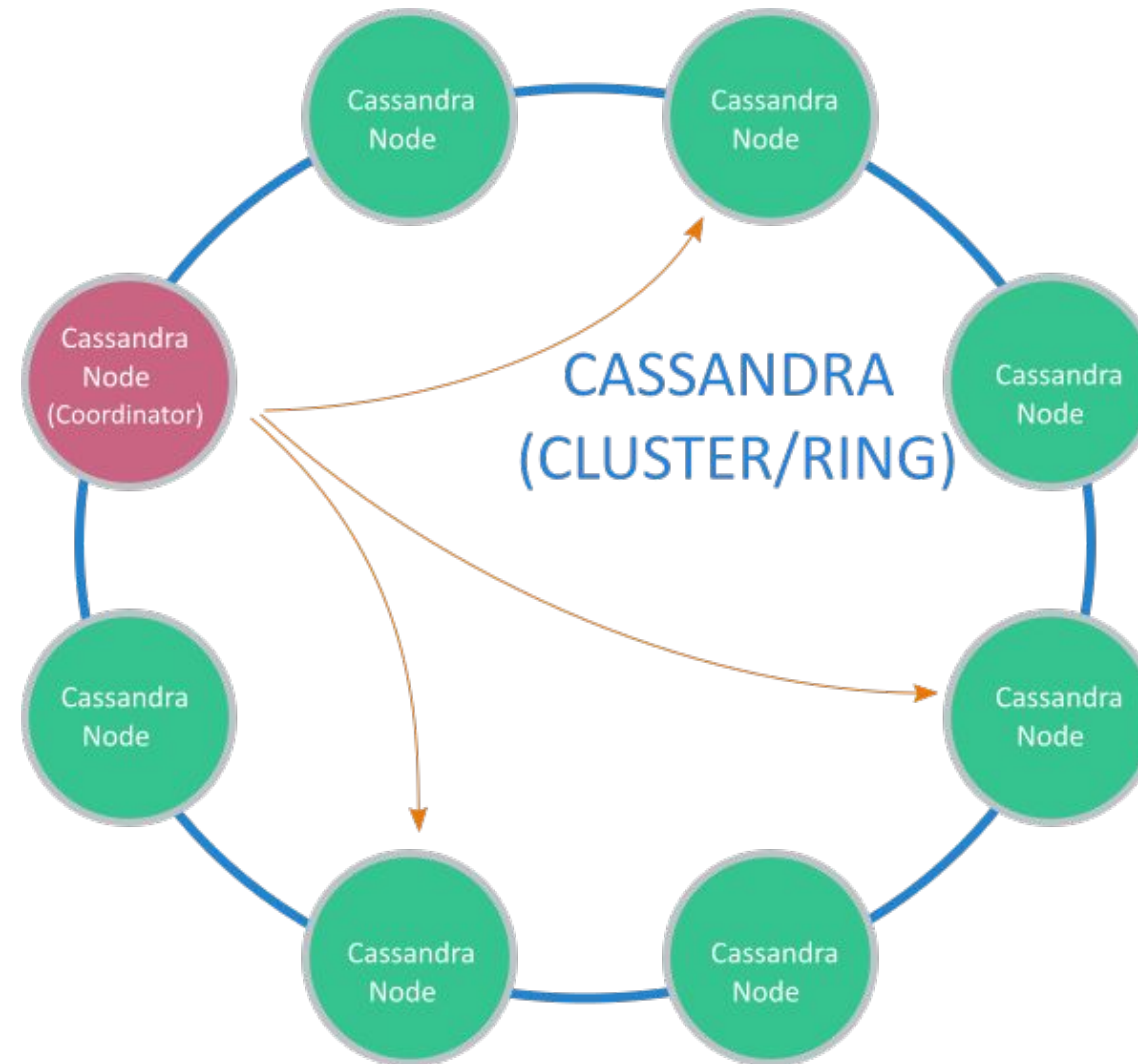
# Stateful Applications



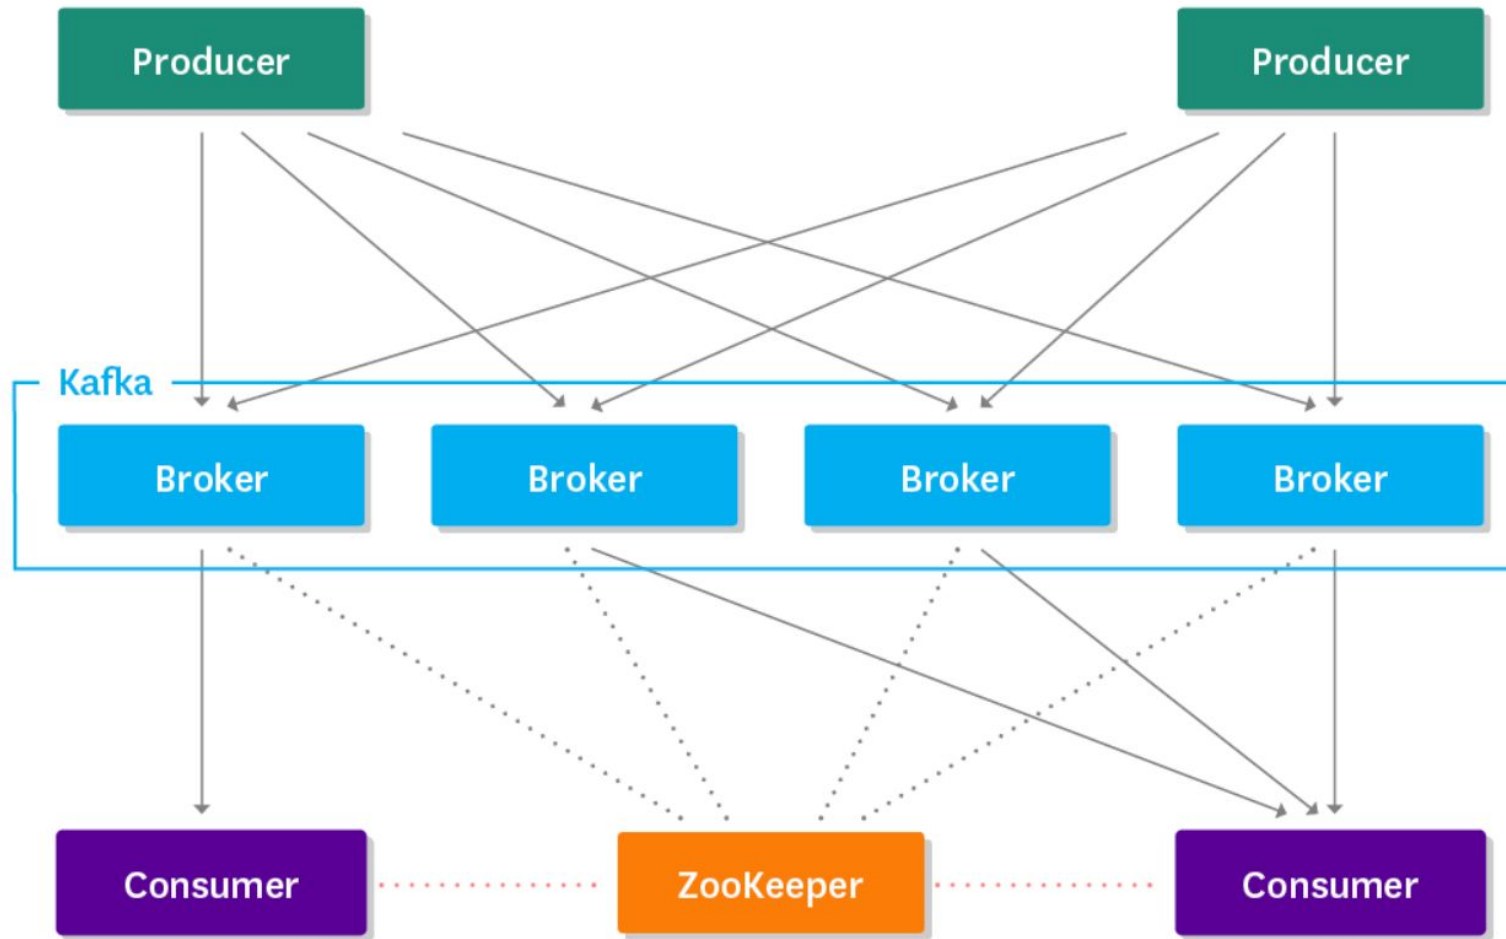
# Stateful Applications



# ***Distributed*** Stateful Applications



# *Distributed* Stateful Applications





# kubernetes

- Focused initially for purely stateless workloads
- Scheduler can move pods around

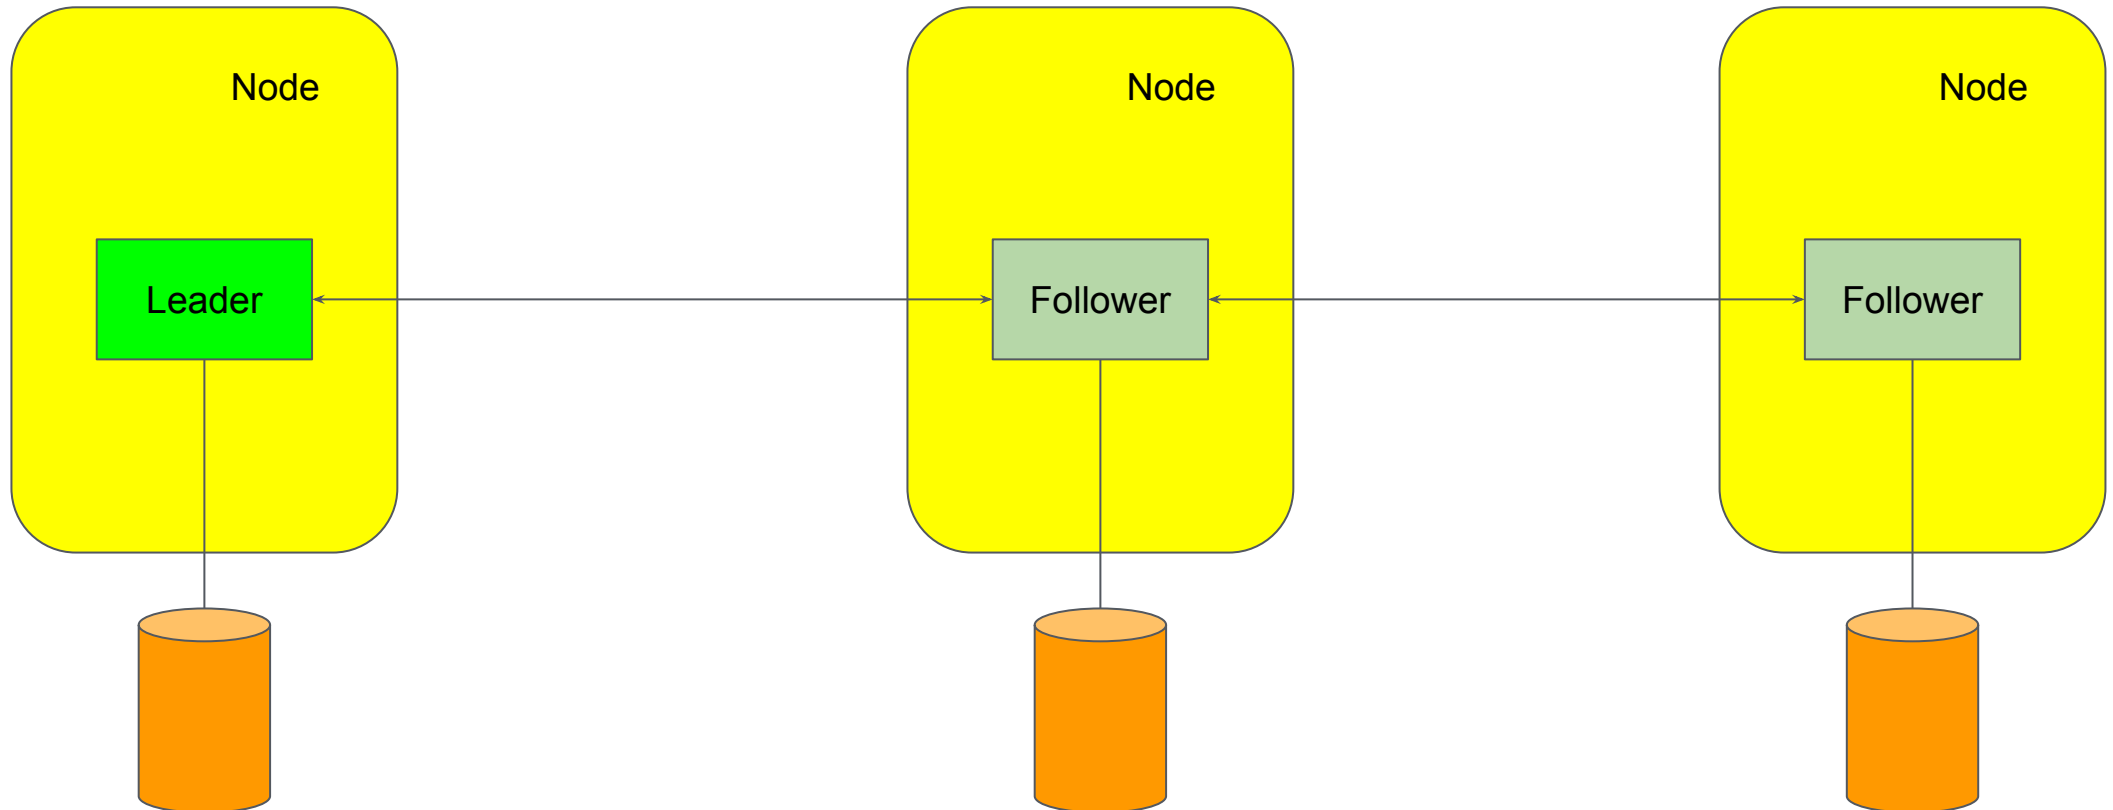


# kubernetes

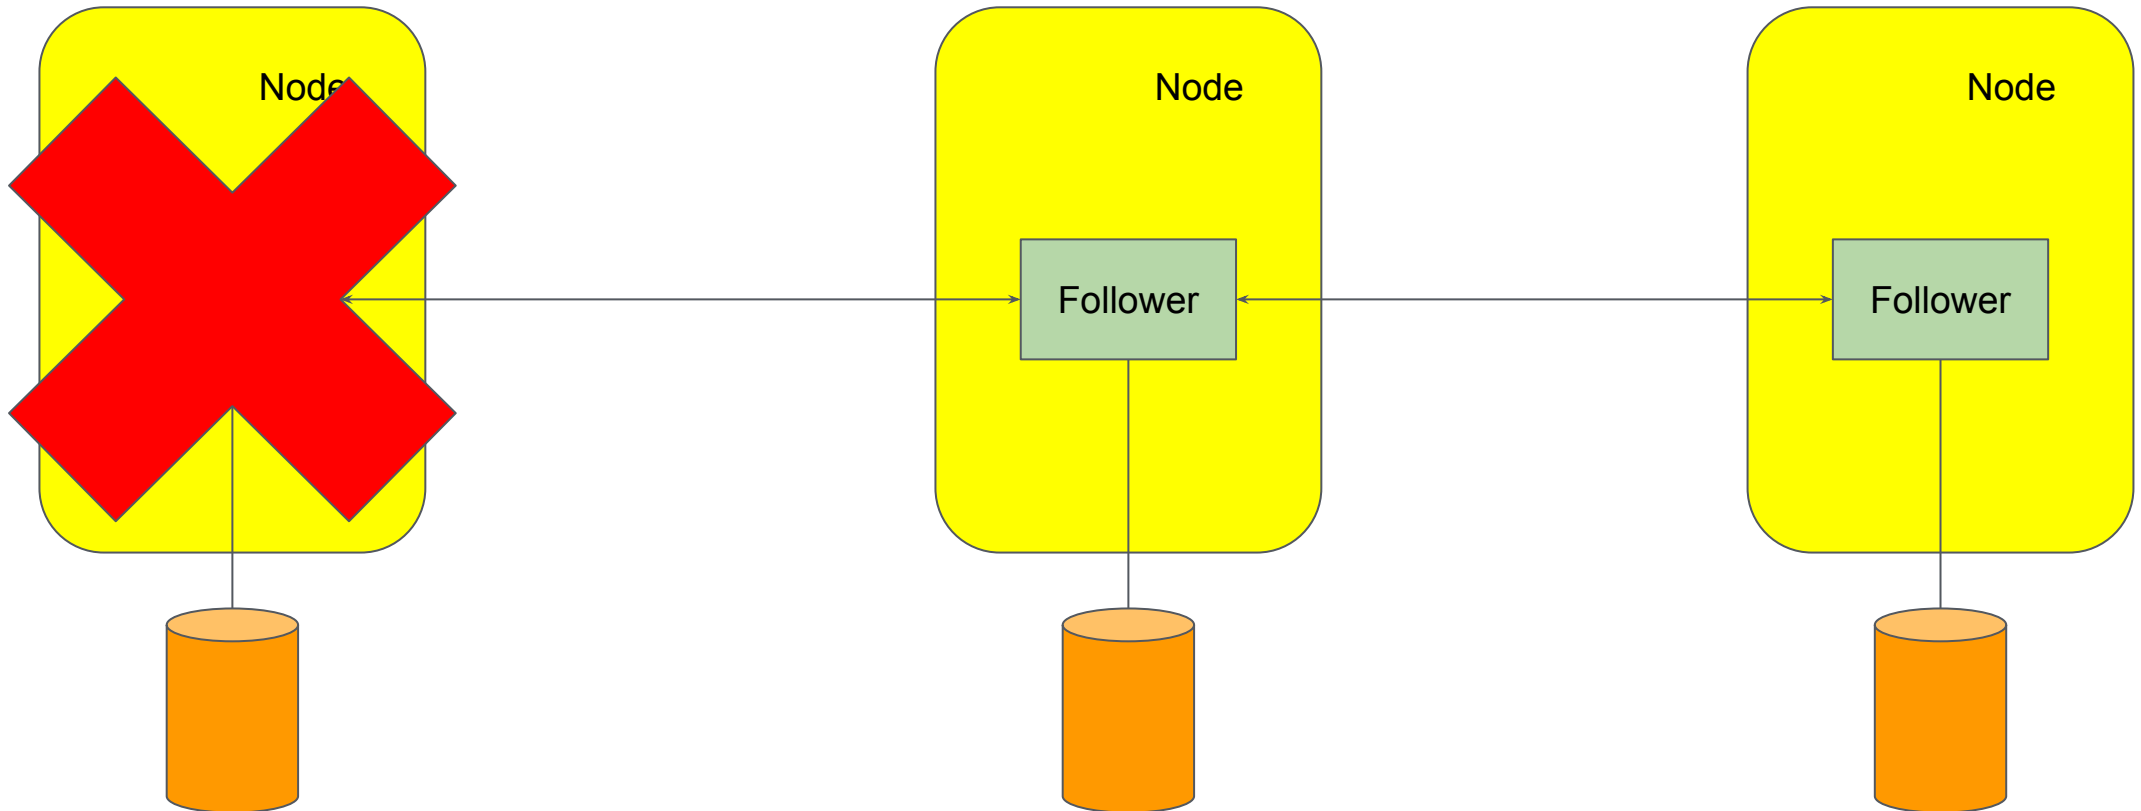
StatefulSets are valuable for applications that require one or more of the following:

- Stable, unique network identifiers.
- Stable, persistent storage.
- Ordered, graceful deployment and scaling.
- Ordered, graceful deletion and termination.
- Ordered, automated rolling updates.

# Kubernetes - StatefulSets

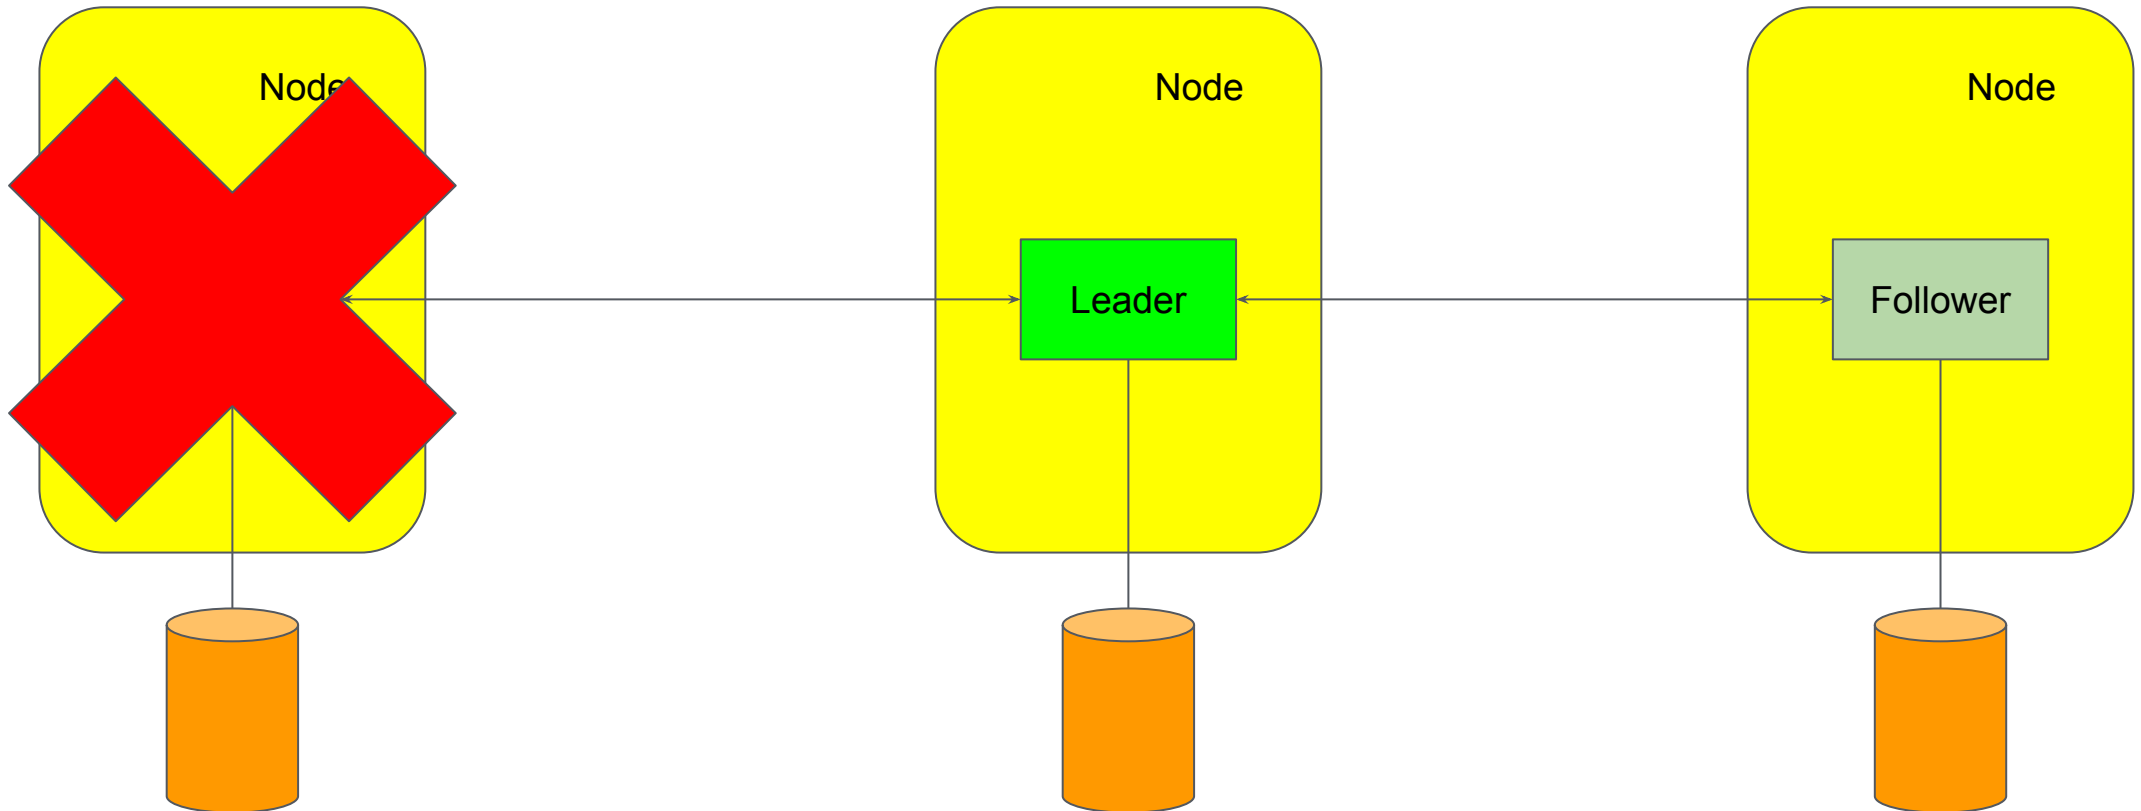


# Kubernetes - StatefulSets







# Kubernetes - StatefulSets




# Kubernetes - StatefulSets



# Kubernetes - StatefulSets



# Kubernetes - StatefulSets





# kubernetes



**Kelsey Hightower** ✓

@kelseyhightower

Following



I'm always going to recommend people exercise extreme caution when running stateful workloads on Kubernetes. Most people who are asking "can I run stateful workloads on Kubernetes" don't have much experience with Kubernetes and often times the workload they are asking about.

3:10 AM - 24 Mar 2019


286 Retweets 901 Likes



# Kubernetes Operators



- Orchestrate stateful applications using K8s API
- Extend API using Custom Resource Definitions
- Encode domain specific operational knowledge
- Upgrades
- Failure and Recovery Scenarios
- Scaling up / down
- Purpose built per application
- “Kubernetes is an Operations API”:
  - <https://blog.atomicinc.com/2018/05/23/kubernetes-is-an-operations-api/>



- Operator manages and monitors lifecycle
- CRD's represent application elements / actions

```
apiVersion:
mysql.presslabs.org/v1alpha1
kind: MysqlCluster
metadata:
  name: my-cluster
spec:
  replicas: 2
  secretName: my-secret
```

```
$ kubectl apply -f mysql-cluster.yaml
```

---

## Operator Framework

- RedHat / IBM project
- Implement using Ansible, Helm charts, or Go
- Existing implementations often don't cover the entire lifecycle
- Ansible and Helm are limited. Go requires 1,000s of lines of controller code

---

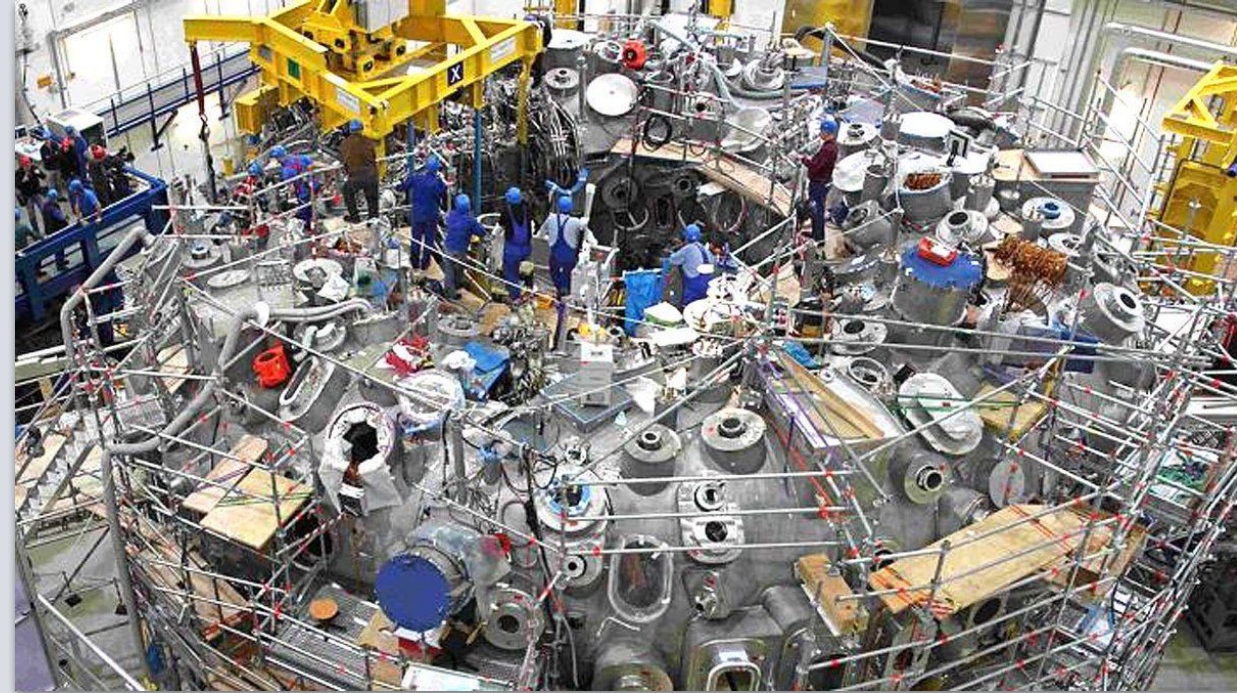
## Kubebuilder

- Kubernetes SIG API Machinery sub-project
- Operators written in Go with a focus on code generation
- Existing implementations often don't cover the entire lifecycle



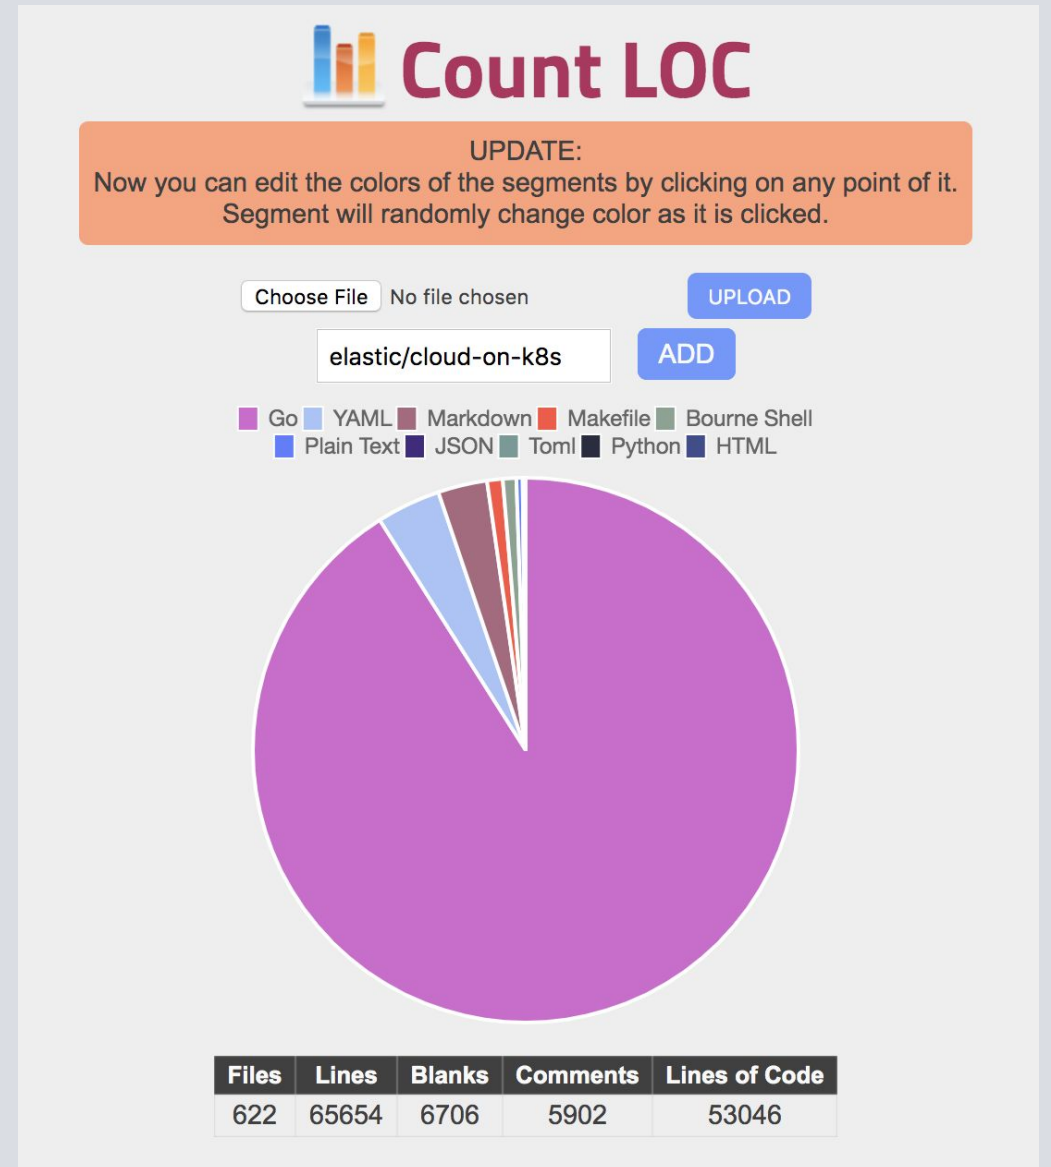
# Developing Operators

- Operators require deep knowledge of Kubernetes internals
- Significant software development undertaking
- May require (10s of) thousands of lines of code
- Controller sprawl can be a thing



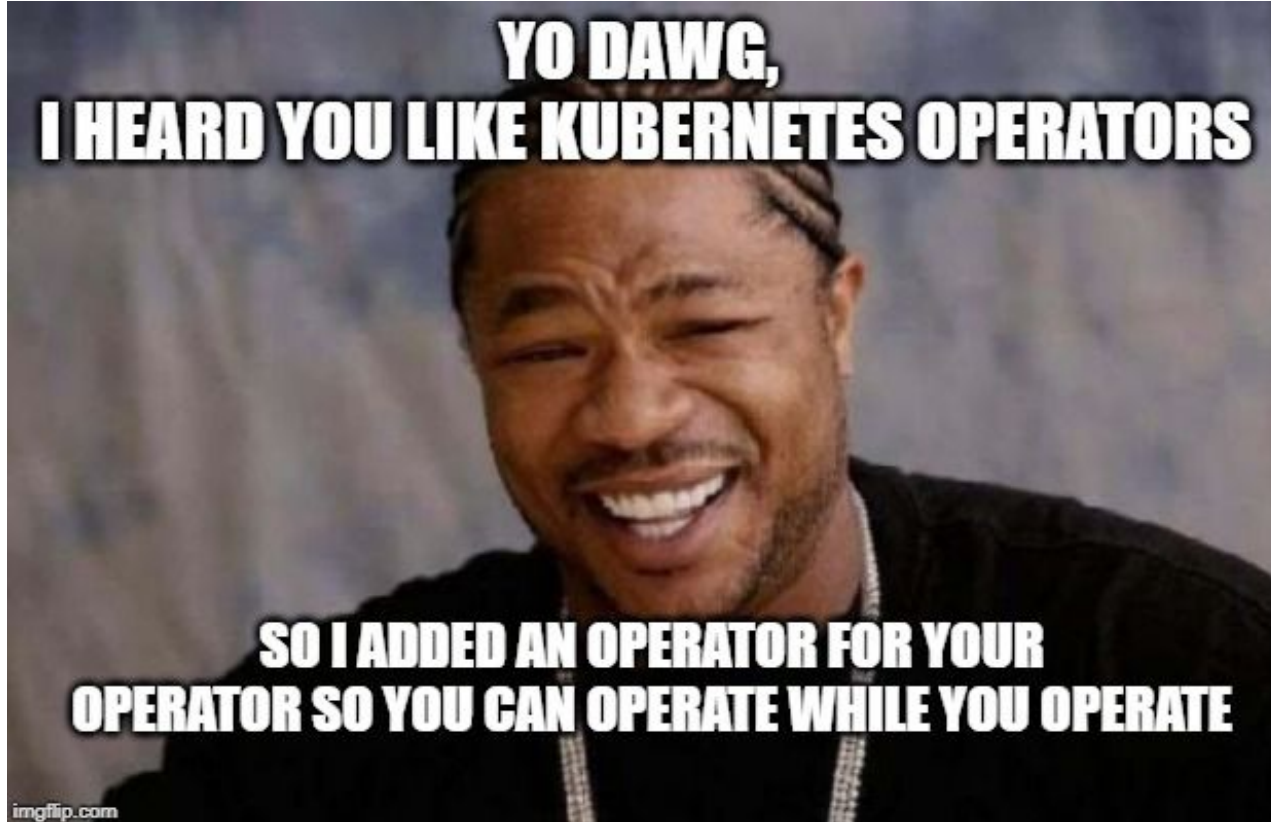
# Developing Operators

- Operators require deep knowledge of Kubernetes internals
- Significant software development undertaking
- May require (10s of) thousands of lines of code
- Controller sprawl can be a thing



👋 operators 👋 operators 👋 operators 👋 operators

operators



**Lachlan Evenson** @LachlanEvenson · 4d

OH: You're going to need an operator to operate that operator

9

2

65



**Andrew Block** @sabre1041 · 4d

and behold the meta-operator

2

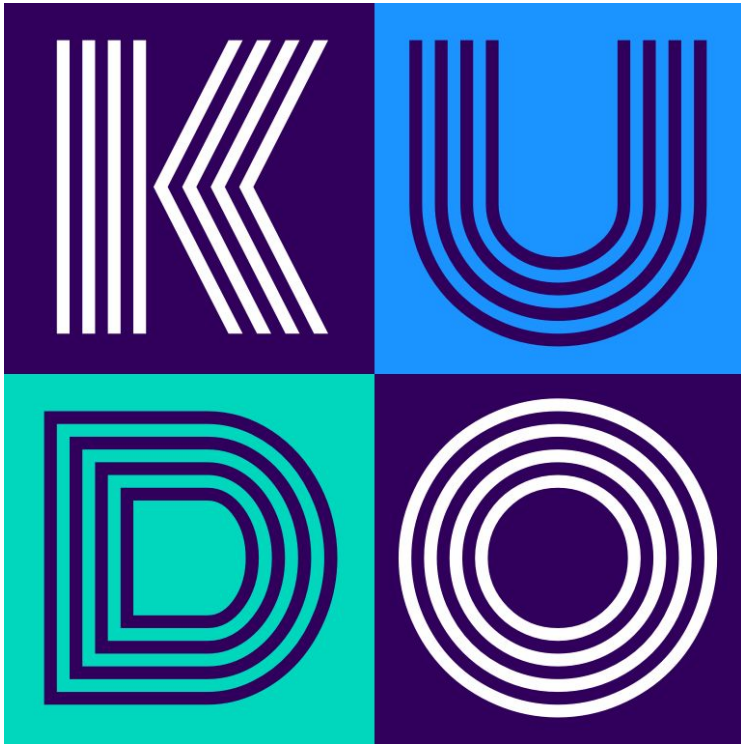
4



**Gaunerd**  
@gaunetes

Replying to @sabre1041 and @LachlanEvenson

Aka KUDO



- Kubernetes Universal Declarative Operator
- A toolkit and runtime for building operators
- Encodes commonality and reuse between lifecycle operations
- Optimised for complex, stateful applications
- Increases developer productivity when **building operators**
- Increases operator productivity when **operating services**
- OS project licensed as Apache 2.0



---

## Operator Framework

- RedHat / IBM project
- Implement using Ansible, Helm charts, or Go
- Existing implementations often don't cover the entire lifecycle
- Ansible and Helm are limited. Go requires 1,000s of lines of controller code

---

## Kubebuilder

- Kubernetes SIG API Machinery sub-project
- Operators written in Go with a focus on code generation
- Existing implementations often don't cover the entire lifecycle

---

## KUDO

- Polymorphic
- Universal Operator
- Built using community projects (Kubebuilder, Kustomize, ...)
- Write Operators as templated YAML manifests
- Provide high level CRDs that represent workloads
- Focused on higher level coordination of software lifecycles
- “Day 2 Operators”

# How KUDO Helps Developers



- Provides abstractions for sequencing lifecycle operations using Kubernetes objects and “plans”, conceptually similar to runbooks
- Encodes commonality and reuse between lifecycle operations
- Reduces boilerplate and code duplication between Operators
- Provides an extension mechanism to create “flavors” of a base Operator for customisation specific to a user’s environment
- Provides ISVs with a tool to ship best practices for Day 2 operations alongside their software
- Ships with testing tool to enable TDD of Kubernetes resources

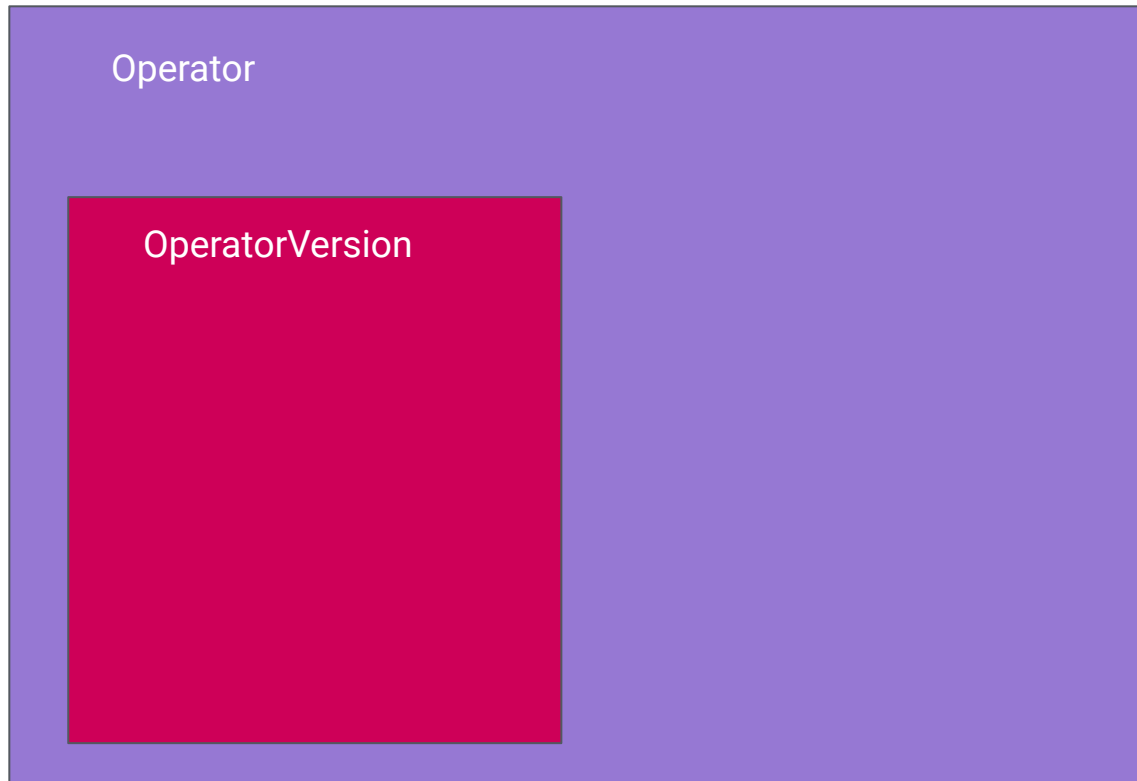
- KUDO provides the `kubectl kudo` plugin to deploy, manage and debug their workloads
  - It's possible to just use kubectl - KUDO is Kubernetes!
- As it's common to deploy multiple Operators to a cluster, KUDO provides a similar API and CLI / workflow experience for all
- All workloads are managed as CRDs, facilitating GitOps
- Existing Operators can be managed by KUDO, natively understanding how to deploy CRDs, custom resource, and other operators, enabling dependencies as part of other workloads
- (Future) Centralised supportability, metrics / alerting, as well as security and RBAC features for Enterprise workloads

A large purple rectangle representing the Operator concept. The word 'Operator' is written in white text in the top-left corner of the rectangle.

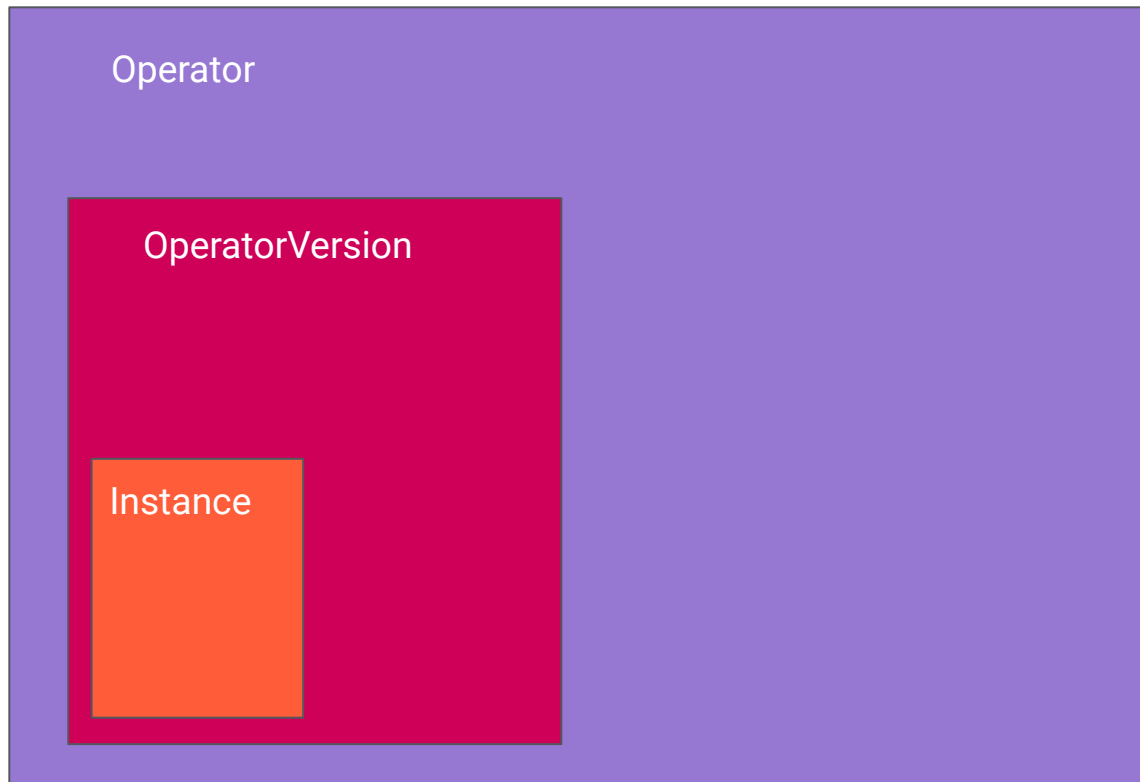
Operator

- High level description of a deployable service
- A deployable service can be anything that you'd want to run on your cluster
- Represented as a CRD object



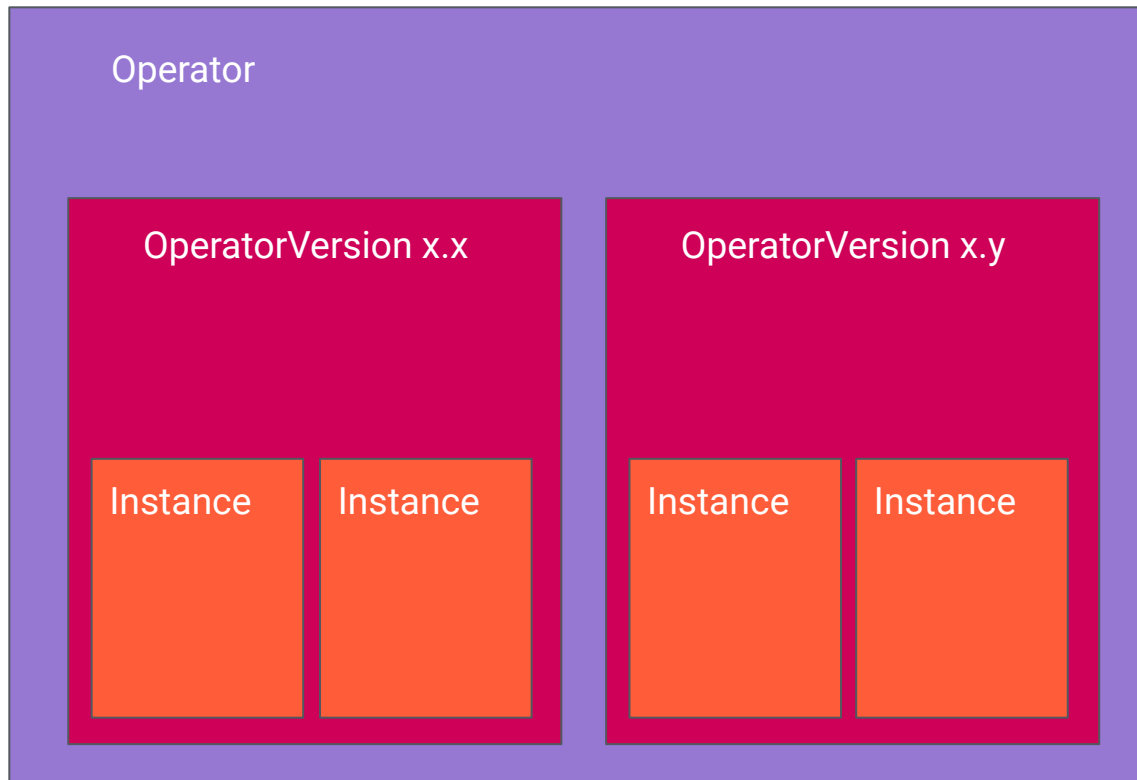


- Implementation of an Operator
- Specific version of a deployable application
- Contains parameters, objects, plans



- Ties application instantiation to an OperatorVersion
- Once created, renders parameters in templates such as services, pods or StatefulSets
- Can create multiple instances of an OperatorVersion within your cluster

# KUDO Concepts - *Instance*



- Ties application instantiation to an OperatorVersion
- Once created, renders parameters in templates such as services, pods or StatefulSets
- Can create multiple instances of an OperatorVersion within your cluster

```
Plan foo
├─ Phase bar
│   ├── Step qux
│   └─ Step quux
└─ Phase baz
    ├── Step quuz
    ├── Step corge
    └─ Step grault
```

- Orchestrate tasks through phases and steps
- A structured 'runbook' which can then be executed by software
- Typically define several plans:
  - Deploy
  - Backup
  - Restore
  - Upgrade
- Phases and steps can be run serial or parallel

- CLI extension to kubectl
- Can still use 'vanilla' kubectl

```
# Install a KUDO package from the official GitHub repo.  
kubectl kudo install <name> [flags]
```

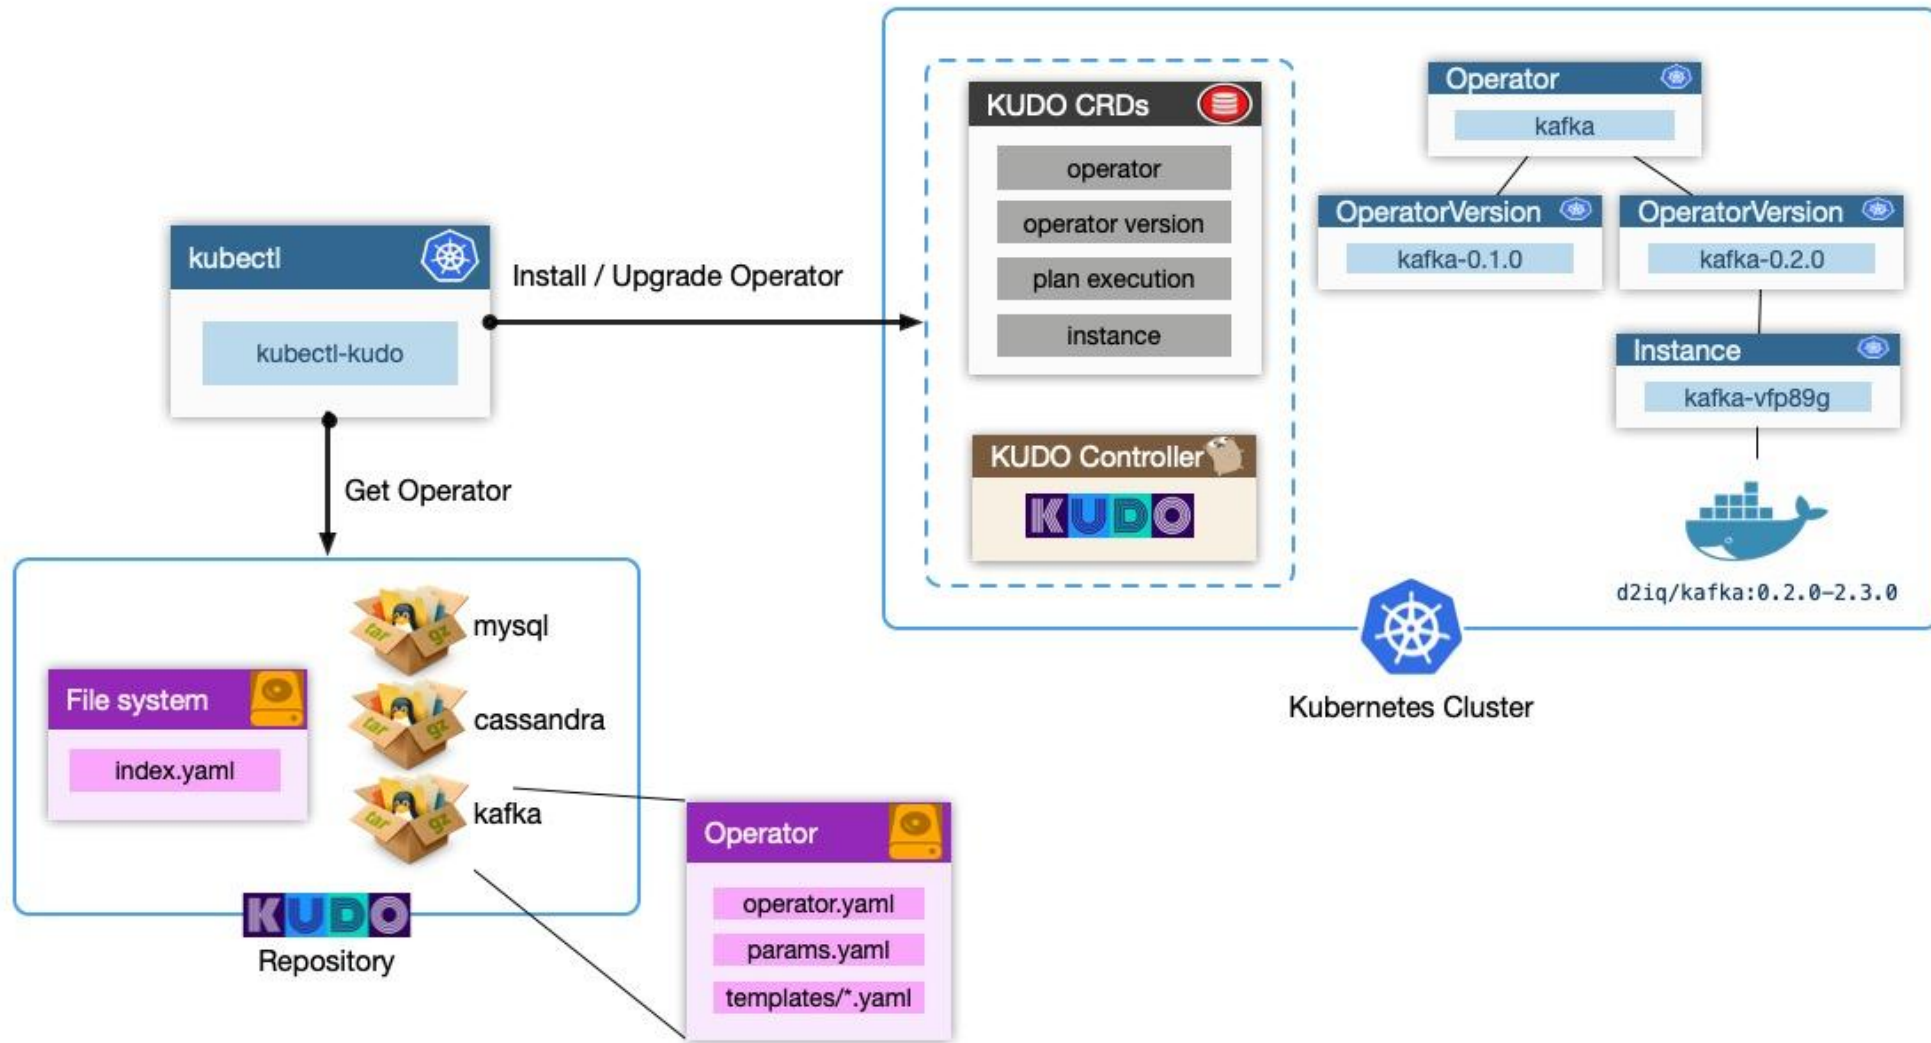
```
# View plan history of a specific package  
kubectl kudo plan history <name> [flags]
```

```
# View all plan history of a specific package  
kubectl kudo plan history [flags]
```

```
# List instances  
kubectl kudo list instances [flags]
```

```
# View plan status  
kubectl kudo plan status [flags]
```

# KUDO Concepts - *Architecture*



Demo - <https://butt.holdings>

D2  
IQ







- Dynamic CRDs
  - Manage the lifecycle of operator CRDs for the operator developers and users
- Operator Dependencies
  - Ability for KUDO to support a wide range of dependencies (from existing instances and connection strings to entirely new dependencies that are KUDO managed), and for tighter control of dependency specification by operator developers.
- Operator Extensions
  - Extend from other formats such as other KUDO operators, Helm charts, or CNAB bundles without forking an operator.
- Something other than YAML! Starlark or CUE likely candidates.
- Pipe Tasks
  - Generation of content which can then be 'piped' to another task
  - E.g certificate generation / creation as part of bootstrap
  - Just landed (<https://github.com/kudobuilder/kudo/pull/1105>) 🎉

- Helm chart
  - Import and extend
- Operator Development
  - Skeleton Generator
  - Linter
  - Snippet / extension library
- KUDO API
- Roadmap here: <https://github.com/orgs/kudobuilder/projects/2>

# KUDO Roadmap - Operator Extensions



Operator Developer Maintained

ACME Corp Maintained

**MySQL**

**MySQL + GKE**


**ACME Corp**

“Standard” infrastructure, plans, CRDs, etc.

Istio, Cloud Storage, GCP Security Rules,  
StackDriver Monitoring, etc.

ACME specific plans. Network policy, special  
operations, cached queries, custom functions,  
etc.

# KUDO Community - Get Involved!



<https://kudo.dev/>



<https://github.com/kudobuilder/kudo>



#kudo <http://slack.k8s.io/>



<https://groups.google.com/forum/#!forum/kudobuilder>



Community Meeting - bi-weekly, Thursdays 10am PT